

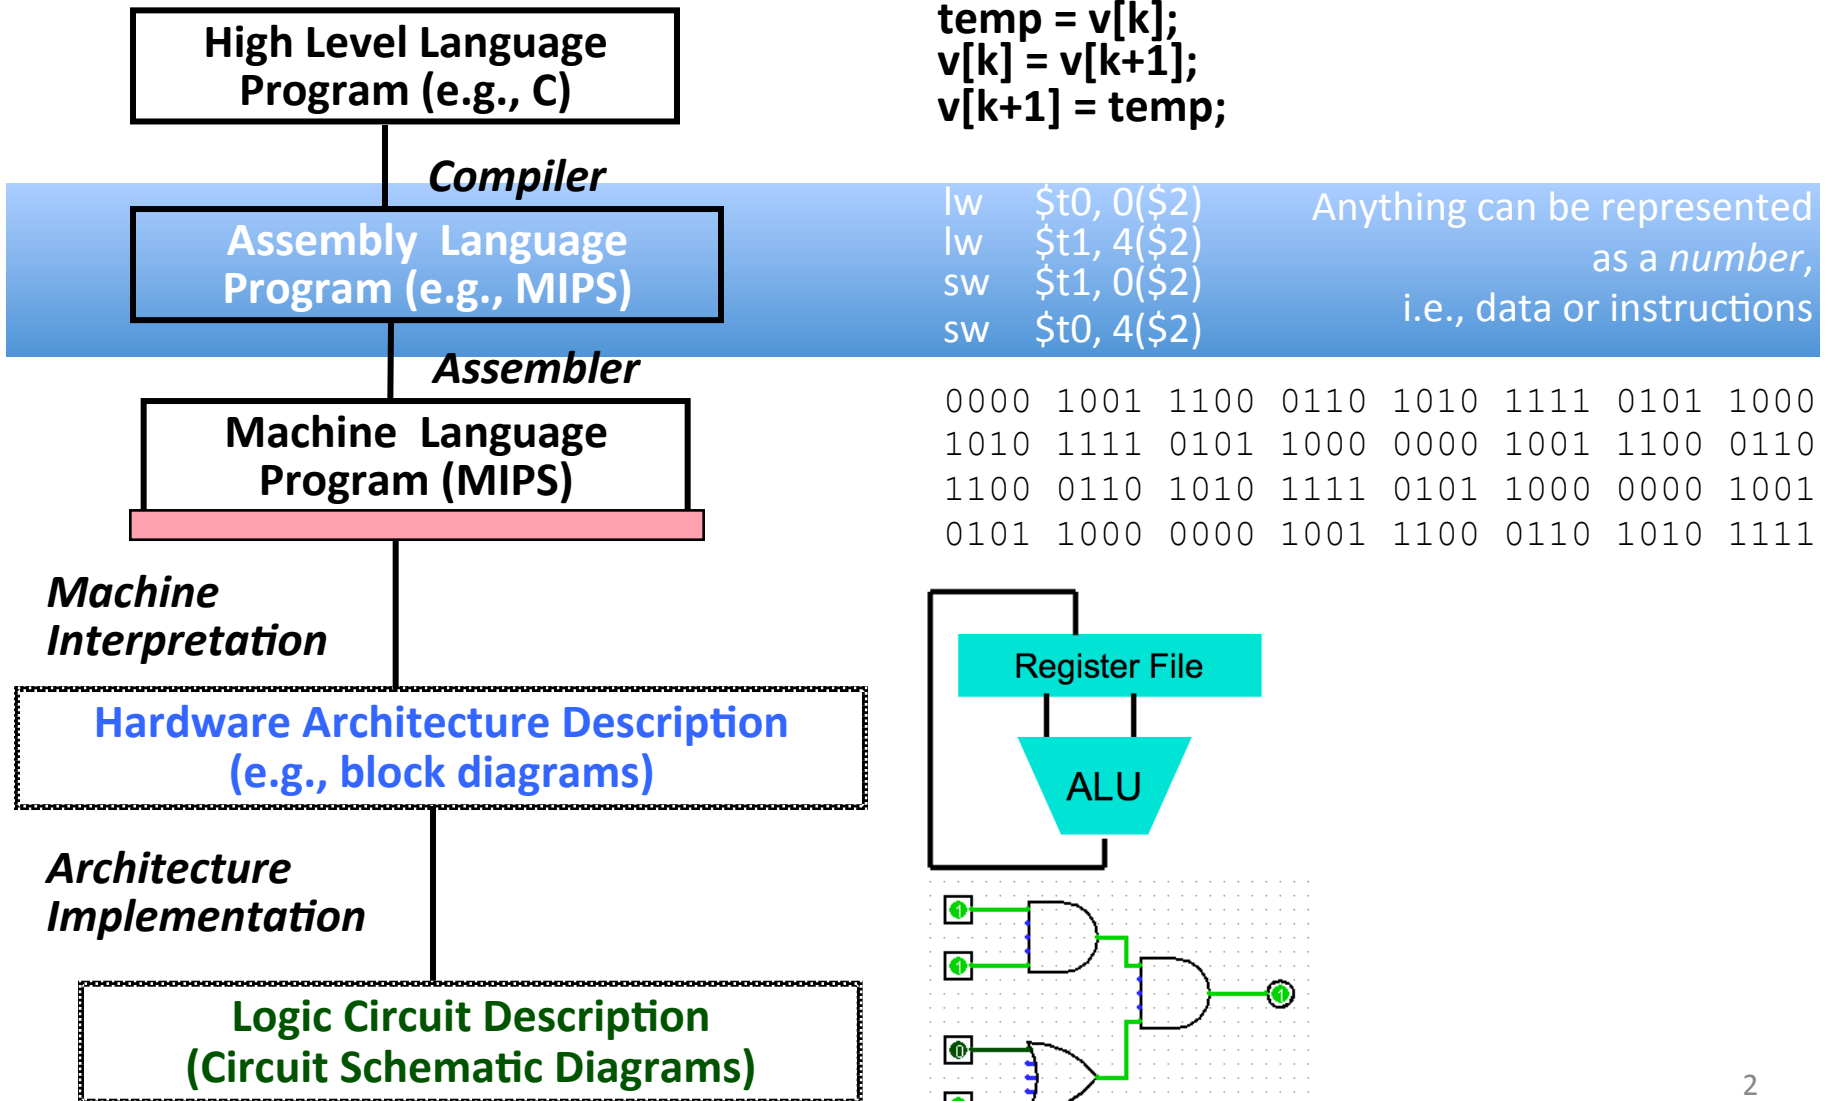
CS 61C: Great Ideas in Computer Architecture

Lecture 5: *Intro to Assembly Language, MIPS Intro*

Instructor: Sagar Karandikar
sagark@eecs.berkeley.edu

<http://inst.eecs.berkeley.edu/~cs61c>

Levels of Representation/ Interpretation



Assembly Language

- Basic job of a CPU: execute lots of *instructions*.
- Instructions are the primitive operations that the CPU may execute.
- Different CPUs implement different sets of instructions. The set of instructions a particular CPU implements is an *Instruction Set Architecture (ISA)*.
 - Examples: ARM, Intel x86, MIPS, RISC-V, IBM/Motorola PowerPC (old Mac), Intel IA64, ...

Instruction Set Architectures

- Early trend was to add more and more instructions to new CPUs to do elaborate operations
 - VAX architecture had an instruction to multiply polynomials!
- RISC philosophy (Cocke IBM, Patterson, Hennessy, 1980s) –
Reduced Instruction Set Computing
 - Keep the instruction set small and simple, makes it easier to build fast hardware.
 - Let software do complicated operations by composing simpler ones.

MIPS Architecture

- MIPS – semiconductor company that built one of the first commercial RISC architectures
- We will study the MIPS architecture in some detail in this class (upper-div arch classes like 150, 152 use a similar ISA, RISC-V)
- Why MIPS instead of Intel x86?
 - MIPS is simple, elegant. Don't want to get bogged down in gritty details.
 - MIPS widely used in embedded apps, x86 little used in embedded, and more embedded computers than PCs

Assembly Variables: Registers

- Unlike HLL like C or Java, assembly cannot use variables
 - Why not? Keep Hardware Simple
- Assembly Operands are registers
 - Limited number of special locations built directly into the hardware
 - Operations can only be performed on these!
- Benefit: Since registers are directly in hardware, they are very fast
(faster than 1 ns - light travels 30cm in 1 ns!!!)

Number of MIPS Registers

- Drawback: Since registers are in hardware, there are a predetermined number of them
 - Solution: MIPS code must be very carefully put together to efficiently use registers
- 32 registers in MIPS
 - Why 32? Smaller is faster, but too small is bad. Goldilocks problem.
- Each MIPS register is 32 bits wide
 - Groups of 32 bits called a word in MIPS

Names of MIPS Registers

- Registers are numbered from 0 to 31
- Each register can be referred to by number or name
- Number references:
 - \$0, \$1, \$2, ... \$30, \$31
- For now:
 - \$16 - \$23 → \$s0 - \$s7 (correspond to C variables)
 - \$8 - \$15 → \$t0 - \$t7 (correspond to temporary variables)
 - Later will explain other 16 register names
- In general, use names to make your code more readable

C, Java variables vs. registers

- In C (and most High Level Languages) variables declared first and given a type
 - Example: `int fahr, celsius;`
`char a, b, c, d, e;`
- Each variable can ONLY represent a value of the type it was declared as (cannot mix and match *int* and *char* variables).
- In Assembly Language, registers have no type; **operation** determines how register contents are treated

Addition and Subtraction of Integers

- Addition in Assembly

- Example: `add $s0, $s1, $s2` (in MIPS)

- Equivalent to: $a = b + c$ (in C)

where C variables \Leftrightarrow MIPS registers are:

$a \Leftrightarrow \$s0, b \Leftrightarrow \$s1, c \Leftrightarrow \$s2$

- Subtraction in Assembly

- Example: `sub $s3, $s4, $s5` (in MIPS)

- Equivalent to: $d = e - f$ (in C)

where C variables \Leftrightarrow MIPS registers are:

$d \Leftrightarrow \$s3, e \Leftrightarrow \$s4, f \Leftrightarrow \$s5$

Addition and Subtraction of Integers

Example 1

- How to do the following C statement?

`a = b + c + d - e;`

- Break into multiple instructions

```
add $t0, $s1, $s2 # temp = b + c
```

```
add $t0, $t0, $s3 # temp = temp + d
```

```
sub $s0, $t0, $s4 # a = temp - e
```

- A single line of C may break up into several lines of MIPS.
- Notice the use of temporary registers – don't want to modify the variable registers \$s
- Everything after the hash mark on each line is ignored (comments)

Immediates

- Immediates are numerical constants.
- They appear often in code, so there are special instructions for them.

- Add Immediate:

`addi $s0,$s1,-10` (in MIPS)

`f = g - 10` (in C)

where MIPS registers `$s0`, `$s1` are associated with C variables **f**, **g**

- Syntax similar to add instruction, except that last argument is a number instead of a register.

`add $s0,$s1,$zero` (in MIPS)

`f = g` (in C)

Overflow in Arithmetic

- Reminder: Overflow occurs when there is a “mistake” in arithmetic due to the limited precision in computers.
- Example (4-bit unsigned numbers):

$$\begin{array}{r} 15 \\ + 3 \\ \hline 18 \end{array}$$

$$\begin{array}{r} 1111 \\ + 0011 \\ \hline 10010 \end{array}$$

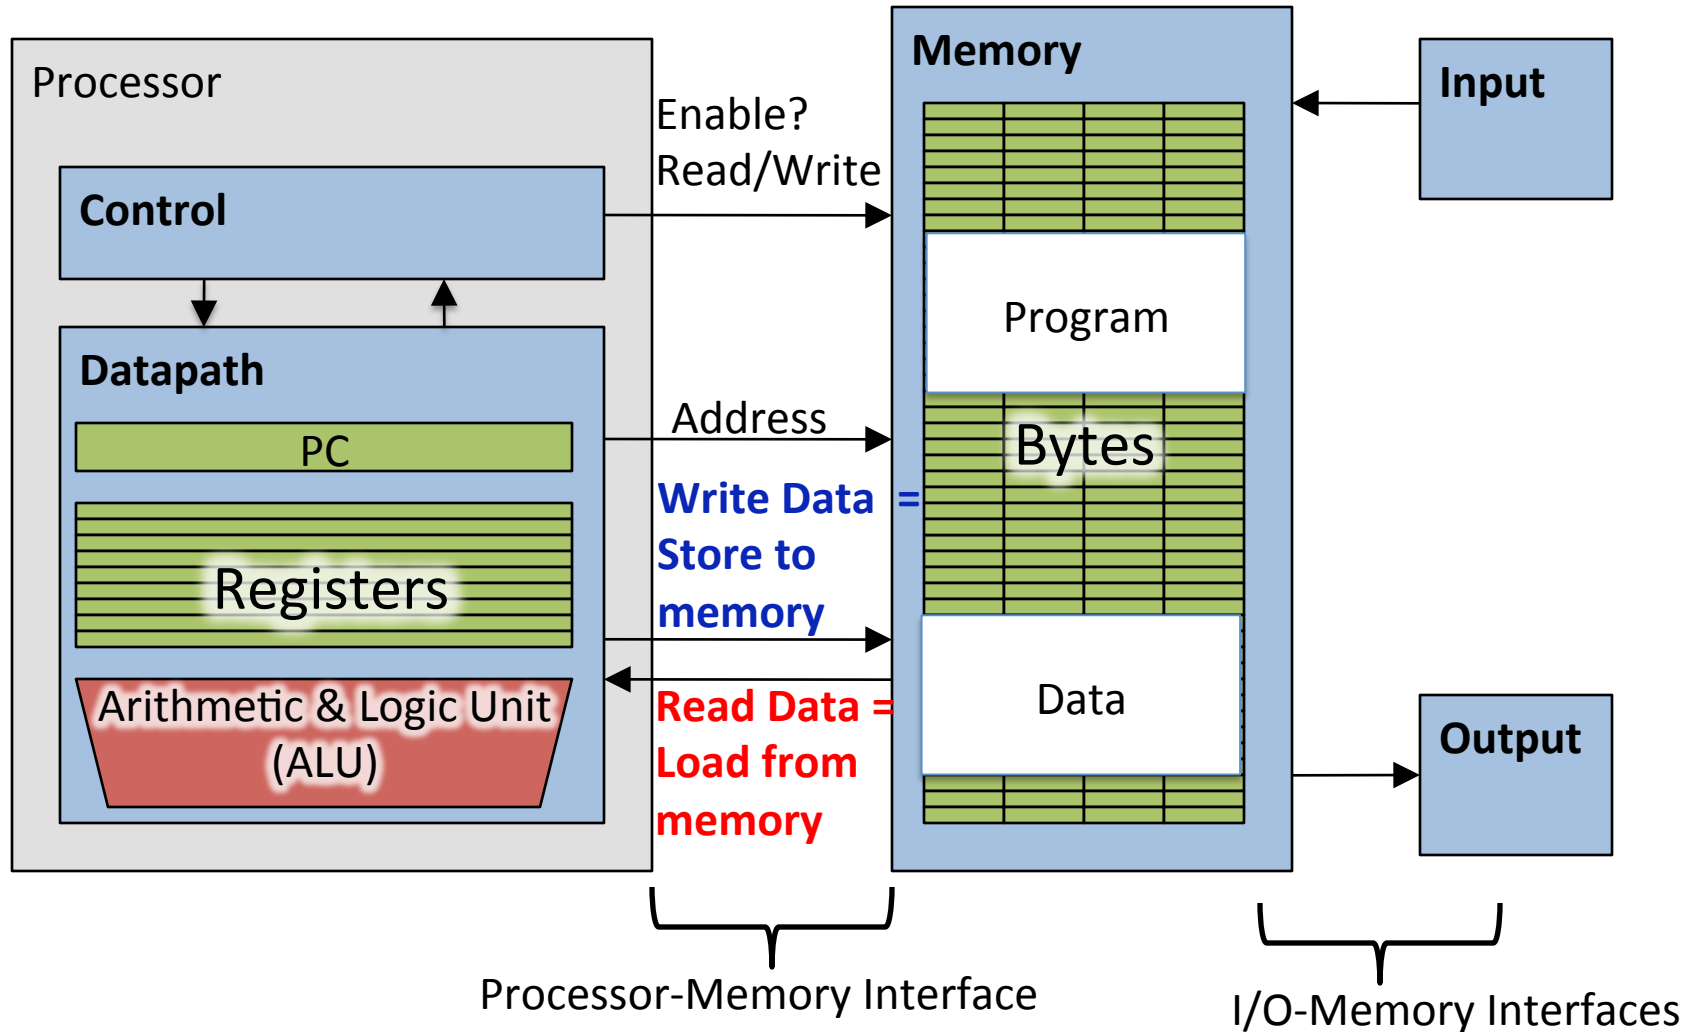
- But we don't have room for 5-bit solution, so the solution would be 0010, which is +2, and “wrong”.

Overflow handling in MIPS

- Some languages detect overflow (Ada), some don't (most C implementations)
- MIPS solution is 2 kinds of arithmetic instructions:
 - These cause overflow to be detected
 - add (**add**)
 - add immediate (**addi**)
 - subtract (**sub**)
 - These do not cause overflow detection
 - add unsigned (**addu**)
 - add immediate unsigned (**addiu**)
 - subtract unsigned (**subu**)
- Compiler selects appropriate arithmetic
 - MIPS C compilers produce **addu, addiu, subu**

Data Transfer:

Load from and Store to memory



Memory Addresses are in Bytes

- Lots of data is smaller than 32 bits, but rarely smaller than 8 bits – works fine if everything is a multiple of 8 bits
- 8 bit chunk is called a *byte*
(1 word = 4 bytes)
- Memory addresses are really in *bytes*, not words
- Word addresses are 4 bytes apart
 - Word address is same as address of leftmost byte (i.e. Big-endian)

Addr of lowest byte in word is addr of word



...
<u>12</u>	13	14	15
<u>8</u>	9	10	11
<u>4</u>	5	6	7
<u>0</u>	1	2	3

Transfer from Memory to Register

- C code

```
int  A[100];  
g = h + A[3];
```

- Using Load Word (lw) in MIPS:

```
lw  $t0, 12($s3)    # Temp reg $t0 gets A[3]  
add $s1, $s2, $t0   # g = h + A[3]
```

Note: \$s3 – base register (pointer)

 12 – offset in bytes

Offset must be a constant known at assembly time

Transfer from Register to Memory

- C code

```
int  A[100];  
A[10] = h + A[3];
```

- Using Store Word (sw) in MIPS:

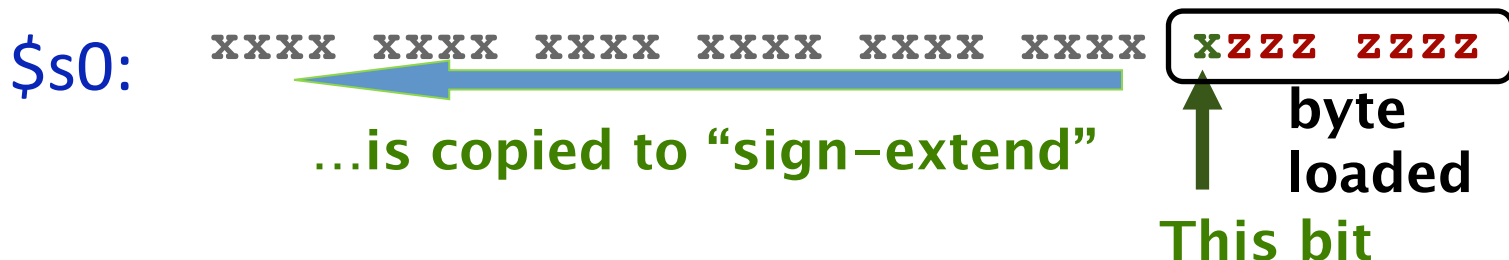
```
lw    $t0, 12($s3)    # Temp reg $t0 gets A[3]  
add   $t0, $s2, $t0    # Temp reg $t0 gets h + A[3]  
sw    $t0, 40($s3)    # A[10] = h + A[3]
```

Note: \$s3 – base register (pointer)
 12, 40 – offsets in bytes

\$s3+12 and \$s3+40 must be multiples of 4

Loading and Storing bytes

- In addition to word data transfers (lw, sw), MIPS has **byte** data transfers:
 - load byte: **lb**
 - store byte: **sb**
- Same format as lw, sw
- E.g., **lb \$s0, 3(\$s1)**
 - contents of memory location with address = sum of “3” + contents of register \$s1 is copied to the low byte position of register **\$s0**.



Speed of Registers vs. Memory

- Given that
 - Registers: 32 words (128 Bytes)
 - Memory: Billions of bytes (2 GB to 8 GB on laptop)
- and the RISC principle is...
 - Smaller is faster
- How much faster are registers than memory??
- About 100-500 times faster!
 - in terms of *latency* of one access

How many hours h on Homework 0?



A: $0 \leq h < 5$

B: $5 \leq h < 10$

C: $10 \leq h < 15$

D: $15 \leq h < 20$

E: $20 \leq h$

Clickers/Peer Instruction

We want to translate $*x = *y + 1$ into MIPS
(x , y pointers to ints, stored in: $\$s0$ $\$s1$)

A: addi $\$s0$, $\$s1$, 1

B: lw $\$s0$, 1($\$s1$)
 sw $\$s1$, 0($\$s0$)

C: lw $\$t0$, 0($\$s1$)
 addi $\$t0$, $\$t0$, 1
 sw $\$t0$, 0($\$s0$)

D: sw $\$t0$, 0($\$s1$)
 addi $\$t0$, $\$t0$, 1
 lw $\$t0$, 0($\$s0$)

E: lw $\$s0$, 1($\$t0$)
 sw $\$s1$, 0($\$t0$)

Break

MIPS Logical Instructions

- Useful to operate on fields of bits within a word
 - e.g., characters within a word (8 bits)
- Operations to pack /unpack bits into words
- Called *logical operations*

Logical operations	C operators	Java operators	MIPS instructions
Bit-by-bit AND	&	&	and
Bit-by-bit OR			or
Bit-by-bit NOT	~	~	not
Shift left	<<	<<	sll
Shift right	>>	>>>	srl

Logic Shifting

- Shift Left: `sll $s1, $s2, 2` #`s1=s2<<2`
 - Store in `$s1` the value from `$s2` shifted 2 bits to the left (they fall off end), inserting 0's on right; `<<` in C.

Before: `0000 0002`_{hex}

`0000 0000 0000 0000 0000 0000 0000 0010`_{two}

After: `0000 0008`_{hex}

`0000 0000 0000 0000 0000 0000 0000 1000`_{two}

What arithmetic effect does shift left have?

- Shift Right: `srl` is opposite shift; `>>`

Arithmetic Shifting

- Shift right arithmetic moves n bits to the right (insert high order sign bit into empty bits)
- For example, if register \$s0 contained
1111 1111 1111 1111 1111 1111 1110 0111_{two} = -25_{ten}
- If executed sra \$s0, \$s0, 4, result is:
1111 1111 1111 1111 1111 1111 1111 1110_{two} = -2_{ten}
- Unfortunately, this is NOT same as dividing by 2^n
 - Fails for odd negative numbers
 - C arithmetic semantics is that division should round towards 0

Computer Decision Making

- Based on computation, do something different
- In programming languages: *if*-statement
- MIPS: *if*-statement instruction is
 `beq register1, register2, L1`
 means: go to statement labeled L1
 if (value in register1) == (value in register2)
 otherwise, go to next statement
- `beq` stands for *branch if equal*
- Other instruction: `bne` for *branch if not equal*

Types of Branches

- **Branch** – change of control flow
- **Conditional Branch** – change control flow depending on outcome of comparison
 - branch *if* equal (`beq`) or branch *if not* equal (`bne`)
- **Unconditional Branch** – always branch
 - a MIPS instruction for this: *jump* (`j`)

Example *if* Statement

- Assuming translations below, compile *if* block

$f \rightarrow \$s0$ $g \rightarrow \$s1$ $h \rightarrow \$s2$

$i \rightarrow \$s3$ $j \rightarrow \$s4$

`if (i == j) bne $s3, $s4, Exit`

`f = g + h; add $s0, $s1, $s2`

`Exit:`

- May need to negate branch condition

Example *if-else* Statement

- Assuming translations below, compile

$f \rightarrow \$s0$ $g \rightarrow \$s1$ $h \rightarrow \$s2$

$i \rightarrow \$s3$ $j \rightarrow \$s4$

if (i == j)

 f = g + h;

else

 f = g - h;

bne \$s3, \$s4, Else

add \$s0, \$s1, \$s2

j Exit

Else: sub \$s0, \$s1, \$s2

Exit:

Administrivia


- Hopefully everyone completed HW0
- HW1 out
- Proj 1 out
 - Make sure you test your code on hive machines, that's where we'll grade them
- First Guerrilla Session this Thursday (07/02) from 5-7pm in the Woz
 - Optional (not part of EPA)
 - Covers Number Rep and MIPS

Administrivia


- Midterm one week from Thursday
 - In this room, at this time
 - One 8.5"x11" handwritten cheatsheet
 - We'll provide a MIPS green sheet
 - No electronics
 - Covers up to and including this Thursday's lecture (07/02)
 - TA-led review session on Monday 07/06 from 5-8pm in HP Auditorium
- Feedback form at the end of lab 2 – tell us how lecture, disc, and lab are going

CS61C In the News

MIPS steers spacecraft to Pluto



- 4 MIPS R3000 32bit CPUs
 - Command and Data handling
 - Guidance and Control



- Launched 2006, first pics in July 2015

<http://www.electronicshweekly.com/news/military-aerospace-electronics/mips-steers-spacecraft-pluto-2015-01/>

CS61C in the News

RISC-V Workshop

The RISC-V Instruction Set Architecture

RISC-V (pronounced "risk-five") is a new instruction set architecture (ISA) that was originally designed to support computer architecture research and education, which we now hope will become a standard open architecture for industry implementations. RISC-V was originally developed in the [Computer Science Division](#) of the EECS Department at the [University of California, Berkeley](#).

CS61C in the News

RISC-V Workshop

- Workshop happening right now at International House
- A modest goal: To become the industry standard ISA (replacing x86, ARM, MIPS, etc.)
- CS150, CS152, CS250, CS252 use RISC-V

CS61C in the News pt. 2

IEEE MILESTONE IN ELECTRICAL ENGINEERING AND COMPUTING

**First RISC (Reduced Instruction-Set Computing) Microprocessor
1980-1982**

UC Berkeley students designed and built the first VLSI reduced instruction-set computer in 1981. The simplified instructions of RISC-I reduced the hardware for instruction decode and control, which enabled a flat 32-bit address space, a large set of registers, and pipelined execution. A good match to C programs and the Unix operating system, RISC-I influenced instruction sets widely used today, including those for game consoles, smartphones and tablets.

February 2015



Inequalities in MIPS

- Until now, we've only tested equalities (`==` and `!=` in C). General programs need to test `<` and `>` as well.

- Introduce MIPS Inequality Instruction:
“Set on Less Than”

Syntax: `slt reg1, reg2, reg3`

Meaning: `if (reg2 < reg3)`
 `reg1 = 1;`
 `else reg1 = 0;`

“set” means “change to 1”,
“reset” means “change to 0”.

Inequalities in MIPS Cont.

- How do we use this? Compile by hand:

if (g < h) goto Less; #g:\$s0, h:\$s1

- Answer: compiled MIPS code...

slt \$t0,\$s0,\$s1 # \$t0 = 1 if g<h

bne \$t0,\$zero,Less # if \$t0!=0 goto Less

- Register \$zero always contains the value 0, so **bne** and **beq** often use it for comparison after an **slt** instruction
- **sltu** treats registers as unsigned

Immediates in Inequalities

- `slti` an immediate version of `slt` to test against constants

`Loop:` . . .

```
slti $t0,$s0,1           # $t0 = 1 if
                          # $s0<1
beq  $t0,$zero,Loop      # goto Loop
                          # if $t0==0
                          # (if ($s0>=1))
```

Loops in C/Assembly

- Simple loop in C; `A[]` is an array of ints

```
do {   g = g + A[i];  
    i = i + j;  
} while (i != h);
```

- Use this mapping:

<code>g</code>	<code>h</code>	<code>i</code>	<code>j</code>	<code>&A[0]</code>
<code>\$s1</code>	<code>\$s2</code>	<code>\$s3</code>	<code>\$s4</code>	<code>\$s5</code>

```
Loop:  sll    $t1, $s3, 2      # $t1 = 4*i  
       addu   $t1, $t1, $s5    # $t1 = addr A+4i  
       lw     $t1, 0($t1)      # $t1 = A[i]  
       add    $s1, $s1, $t1    # g = g + A[i]  
       addu   $s3, $s3, $s4    # i = i + j  
       bne    $s3, $s2, Loop   # goto Loop  
                                   # if i != h
```


Which of the following are true for the `addiu` instruction?



1: The instruction performs a different operation at the hardware level than `add` (excluding overflow reporting)

2: The instruction tells the hardware not to report an overflow

3: The instruction sign extends the immediate

A: F, F, F

B: T, T, T

C: F, T, F

D: F, T, T

E: T, F, F

And In Conclusion ...

- Computer words and vocabulary are called instructions and instruction set respectively
- MIPS is example RISC instruction set in this class
- Rigid format: 1 operation, 2 source operands, 1 destination
 - `add, sub, mul, div, and, or, sll, srl, sra`
 - `lw, sw, lb, sb` to move data to/from registers from/to memory
 - `beq, bne, j, slt, slti` for decision/flow control
- Simple mappings from arithmetic expressions, array access, if-then-else in C to MIPS instructions