

# CS 61C: Great Ideas in Computer Architecture

## Lecture 14: *Caches, Part I*

Instructor: Sagar Karandikar  
sagark@eecs.berkeley.edu

<http://inst.eecs.berkeley.edu/~cs61c>

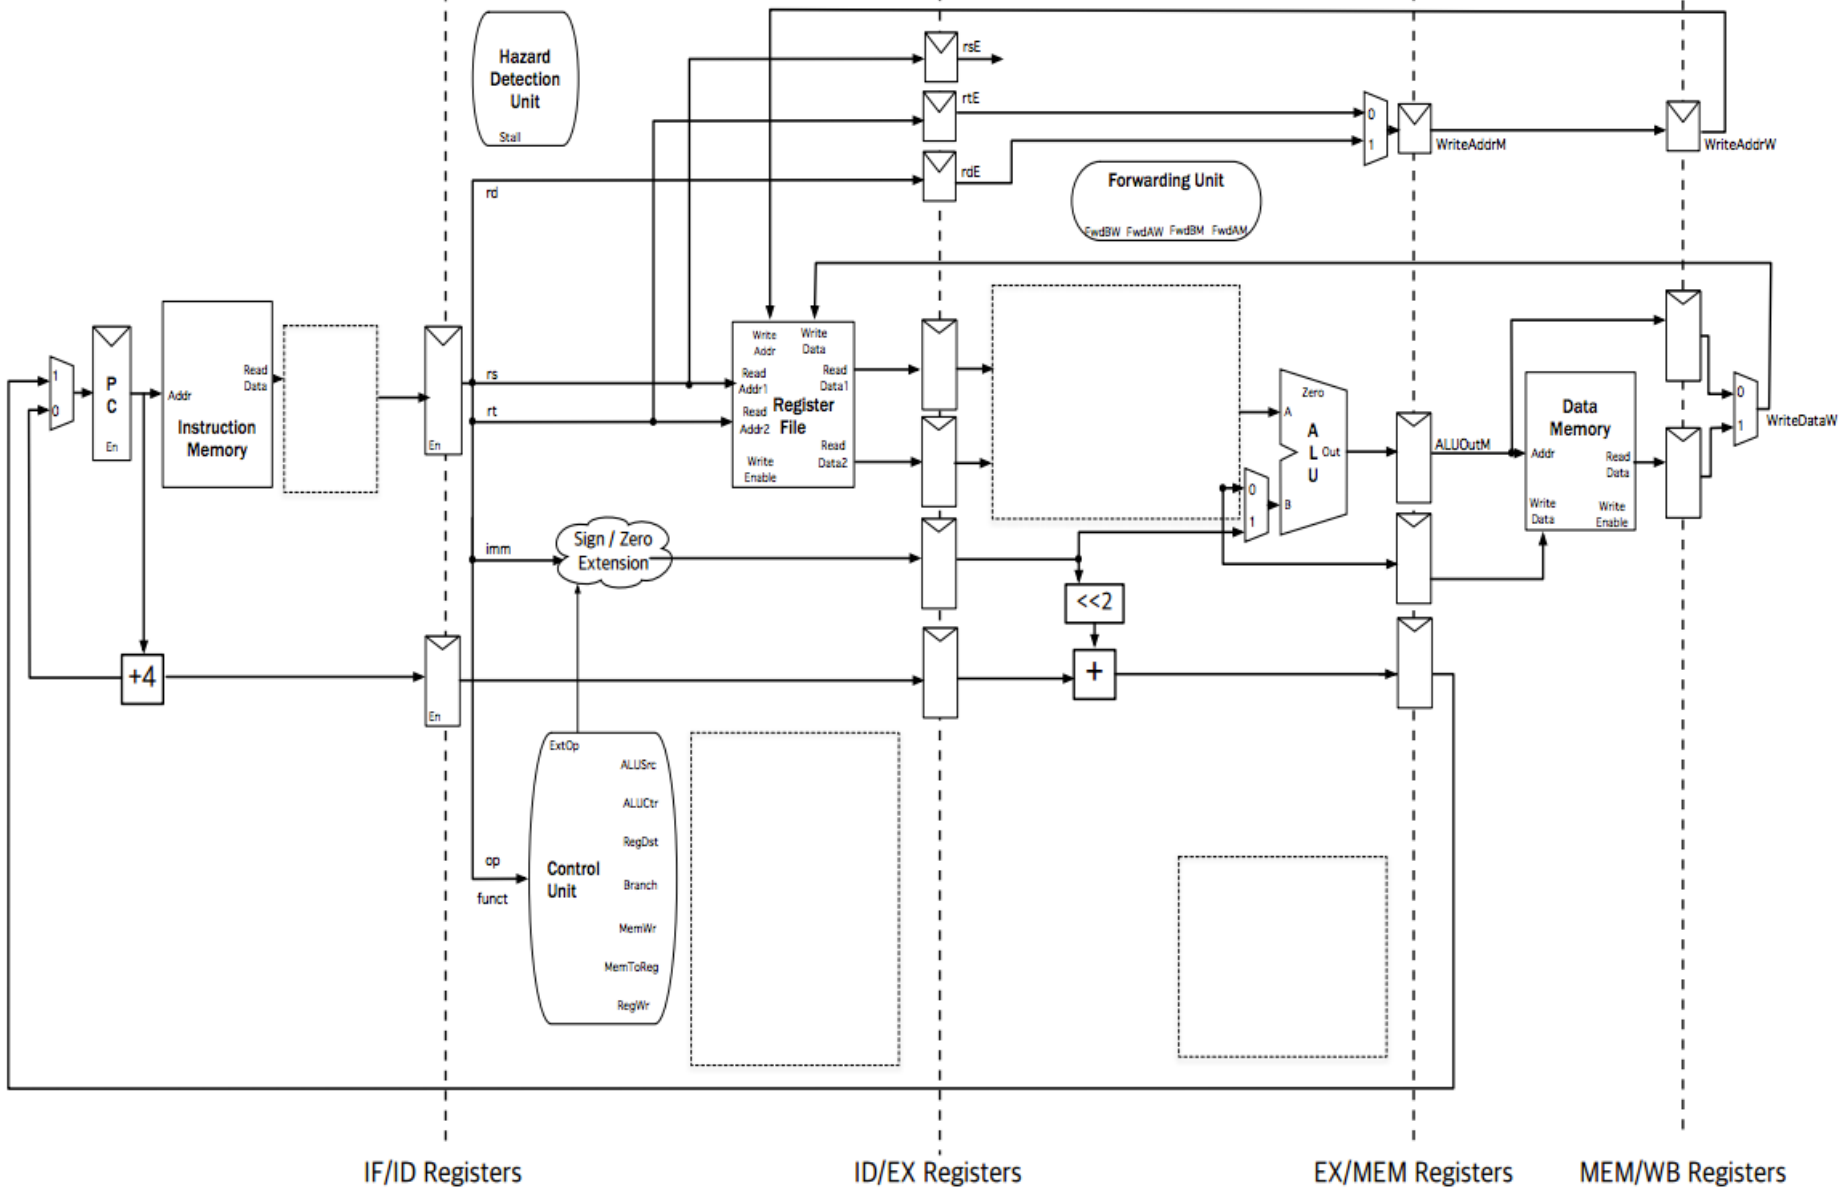
# Review: Single Cycle Performance

- Assume time for actions are
  - 100ps for register read or write; 200ps for other events
- Clock period is?

Instr	Instr fetch	Register read	ALU op	Memory access	Register write	Total time
lw	200ps	100 ps	200ps	200ps	100 ps	800ps
sw	200ps	100 ps	200ps	200ps		700ps
R-format	200ps	100 ps	200ps		100 ps	600ps
beq	200ps	100 ps	200ps			500ps

- What can we do to improve clock rate?
- Will this improve performance as well?
  - Want increased clock rate to mean faster programs

Instruction Fetch      Instruction Decode / Register Read      Execute      Memory      Write Back




# Review: Pipelining Performance

## Single-cycle

$T_c = 800 \text{ ps}$


$f = 1.25\text{GHz}$



## Pipelined

$T_c = 200 \text{ ps}$

$f = 5\text{GHz}$



# Review: Pipelining Hazards

A *hazard* is a situation that prevents starting the next instruction in the next clock cycle

## *1) Structural hazard*

- A required resource is busy  
(e.g. needed in multiple stages)

## *2) Data hazard*

- Data dependency between instructions
- Need to wait for previous instruction to complete its data read/write

## *3) Control hazard*

- Flow of execution depends on previous instruction

# Review: Resolving Control Hazards

- Jumps would require inserting one bubble (since we don't know the jump immediate until decode)
  - Jump delay slot lets us put a useful instruction here
- Branches would require inserting two bubbles (since we don't do the equality comparison until the ALU stage)
  - Move equality comparison to decode stage, add branch delay slot

# Example: Nondelayed vs. Delayed Jump

## Nondelayed Jump

**or \$8, \$9, \$10**

**add \$1, \$2, \$3**

**sub \$4, \$5, \$6**

**jal Foo**

\$ra set to  
addr of xor:  
PC+4

**xor \$10, \$1, \$11**

**Foo:**

## Delayed Jump

**add \$1, \$2, \$3**

**sub \$4, \$5, \$6**

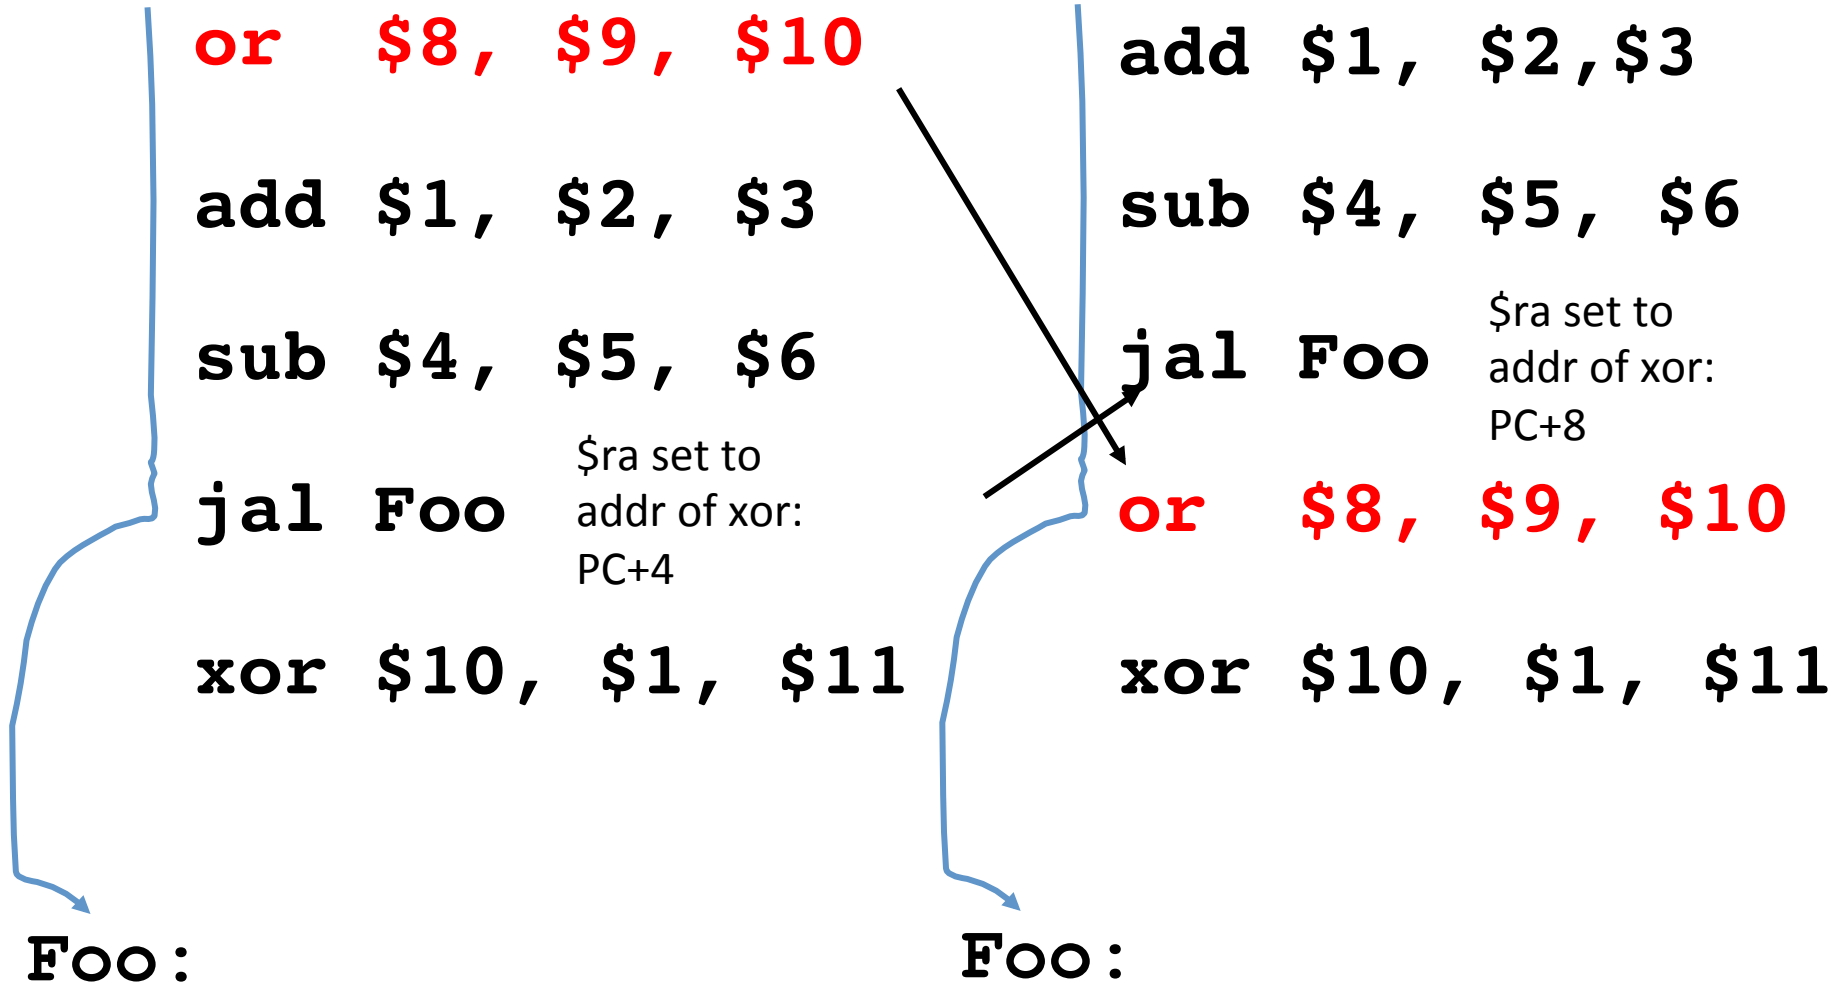
**jal Foo**

\$ra set to  
addr of xor:  
PC+8

**or \$8, \$9, \$10**

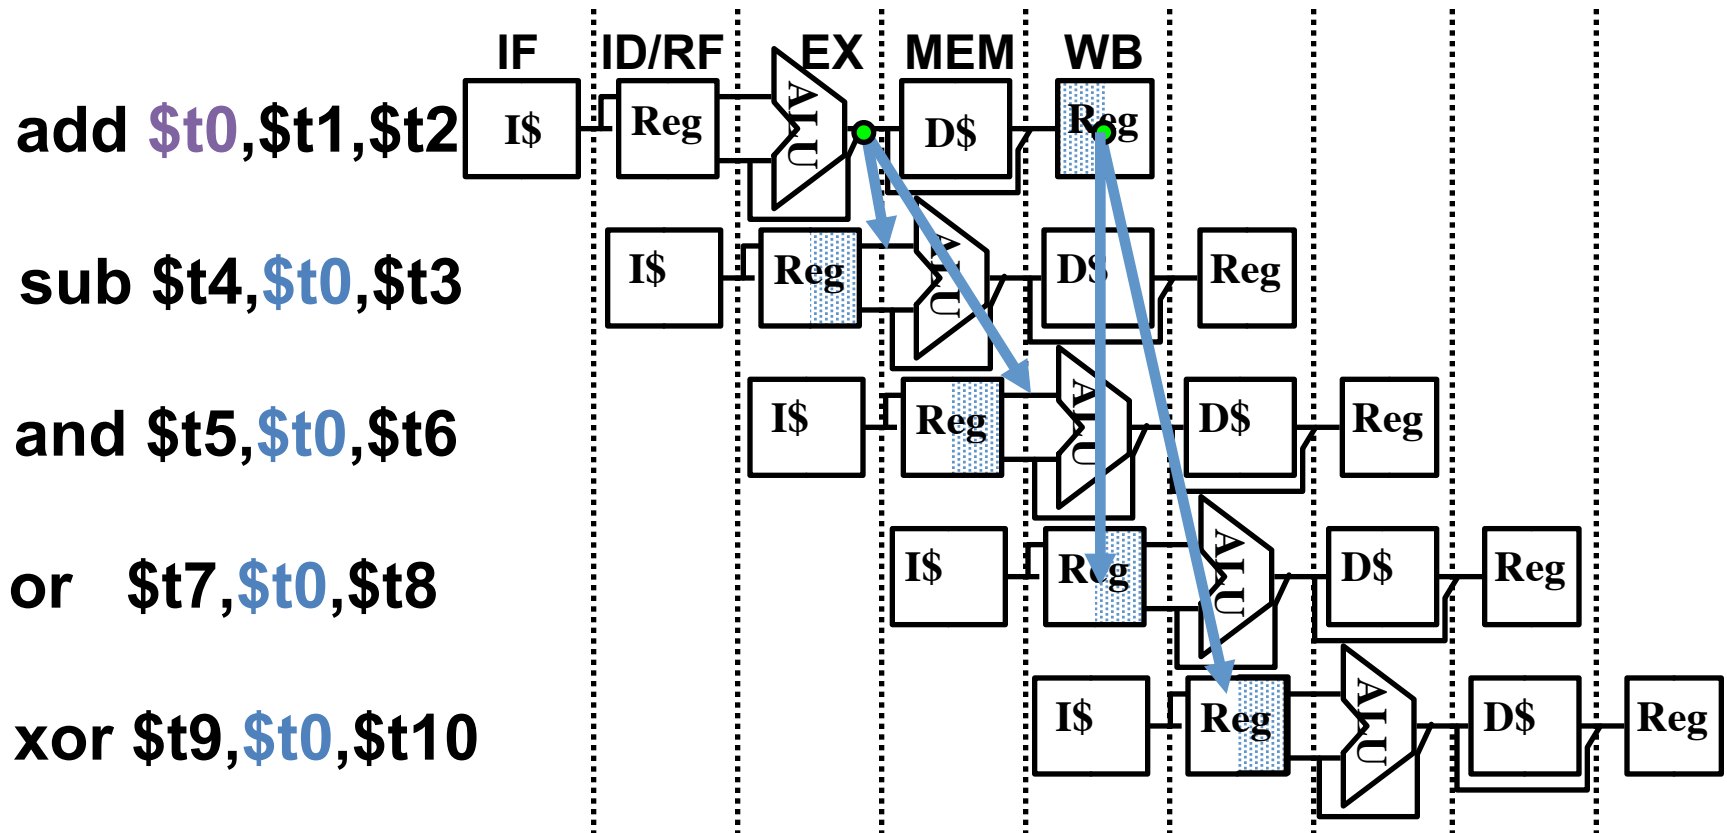
**xor \$10, \$1, \$11**

**Foo:**



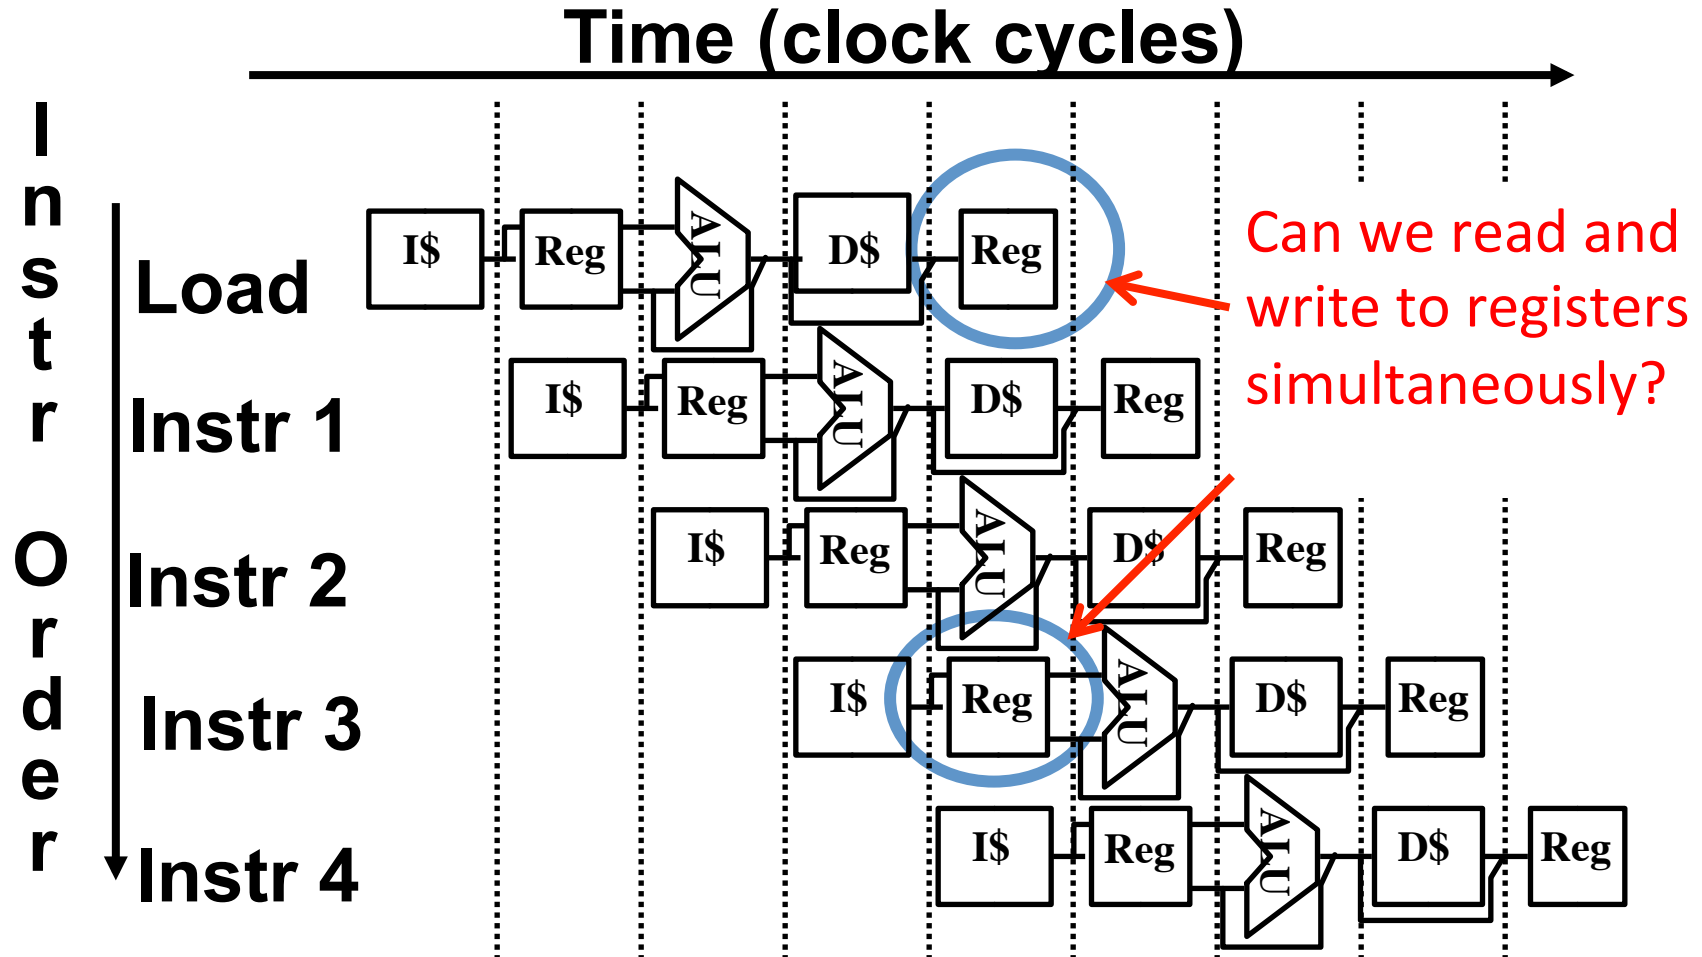
# Review: Resolving (many) Data Hazards: Forwarding

- Forward result as soon as it is available
  - OK that it's not stored in RegFile yet

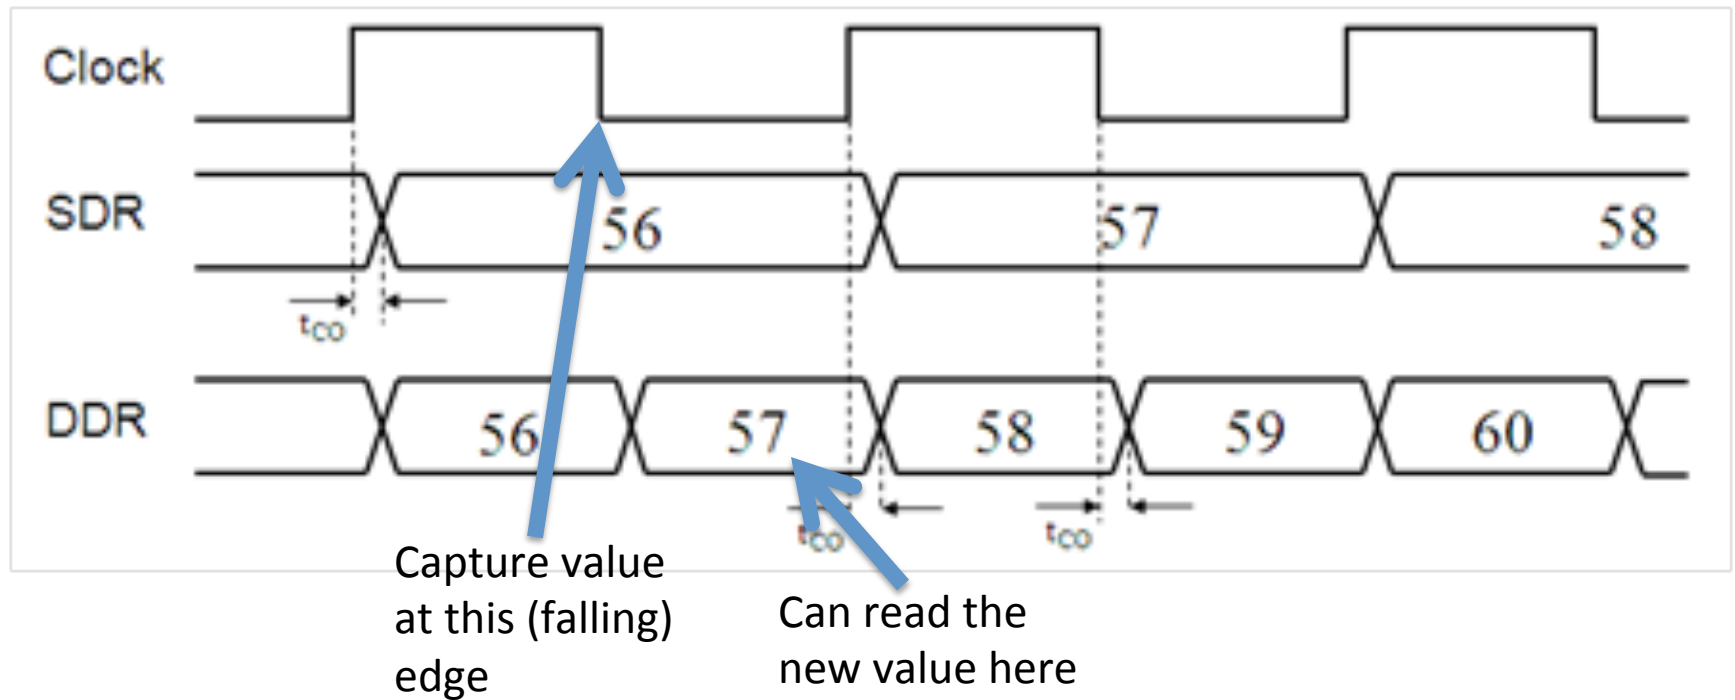




# Review: Structural Hazard #1: Registers



# Review: Solving Regfile Structural Hazard: Double Pumping (2/3)




\* **This is not how registers work in general!**

# Review: Structural Hazard #2: Single Memory

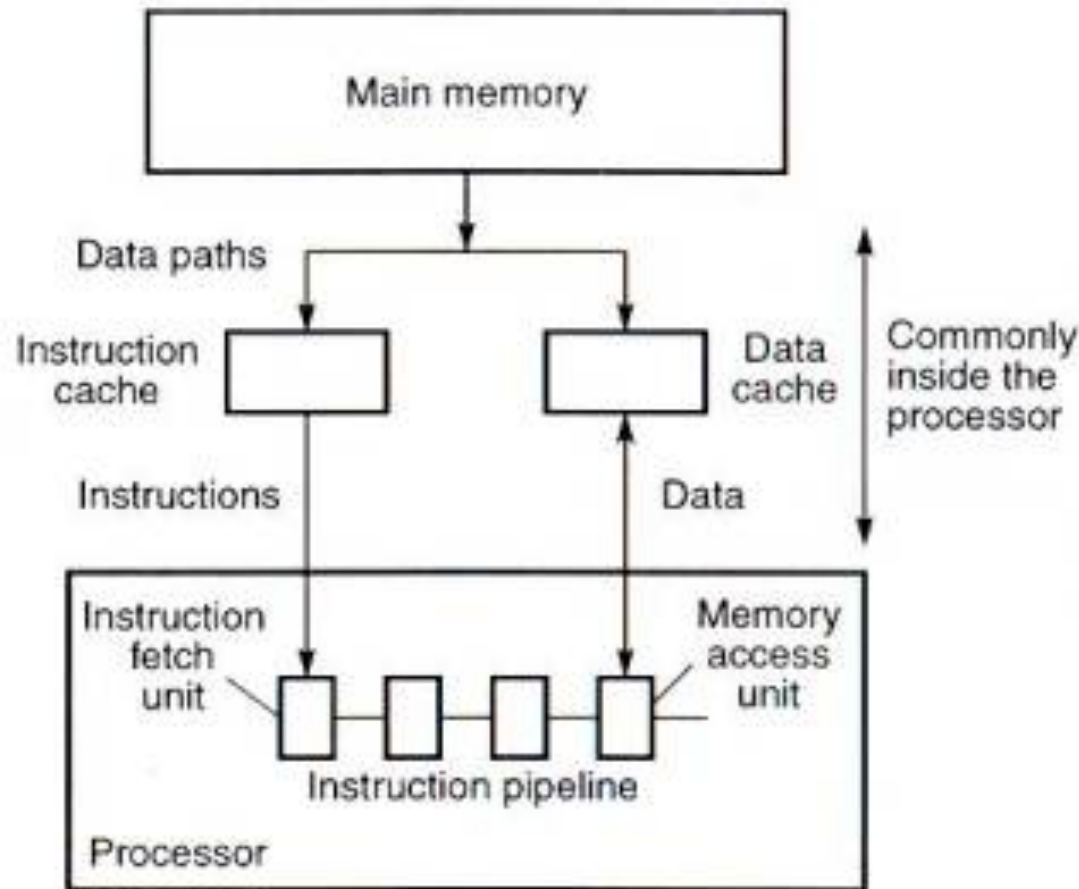
- MIPS pipeline with a single memory?
  - Load/Store requires memory access for data
  - Instruction fetch would have to *stall* for that cycle
    - Causes a pipeline “*bubble*”
- Hence, pipelined datapaths require separate instruction/data memories
  - Separate L1 I\$ and L1 D\$ take care of this

# Review: Structural Hazard #2: Single Memory



# Review: Solving Structural Hazard #2 with Caches

Today,  
we'll add  
caches



# New-School Machine Structures (It's a bit more complicated!)

*Software*

*Hardware*

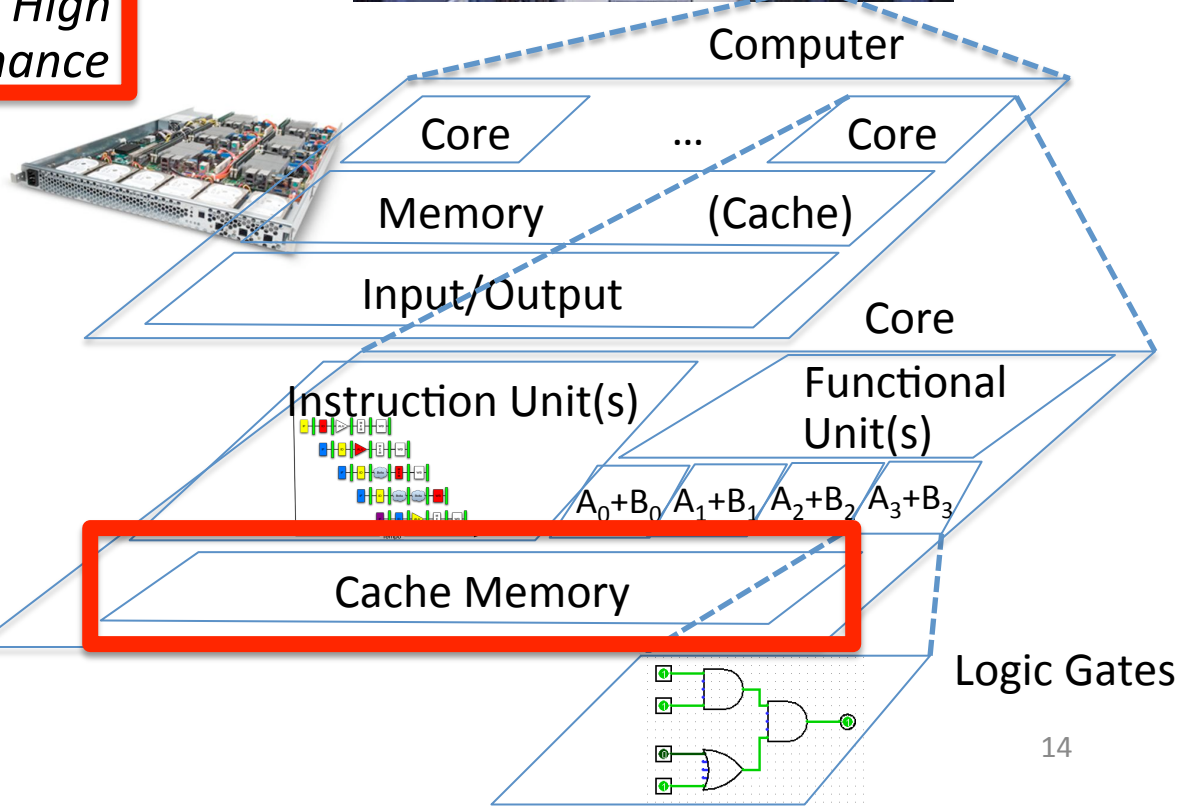
Warehouse  
Scale  
Computer

Smart  
Phone




*Harness  
Parallelism &  
Achieve High  
Performance*

How do  
we know?




- Parallel Requests  
Assigned to computer  
e.g., Search “Katz”
- Parallel Threads  
Assigned to core  
e.g., Lookup, Ads
- Parallel Instructions  
>1 instruction @ one time  
e.g., 5 pipelined instructions
- Parallel Data  
>1 data item @ one time  
e.g., Add of 4 pairs of words
- Hardware descriptions  
All gates @ one time
- Programming Languages

# Components of a Computer



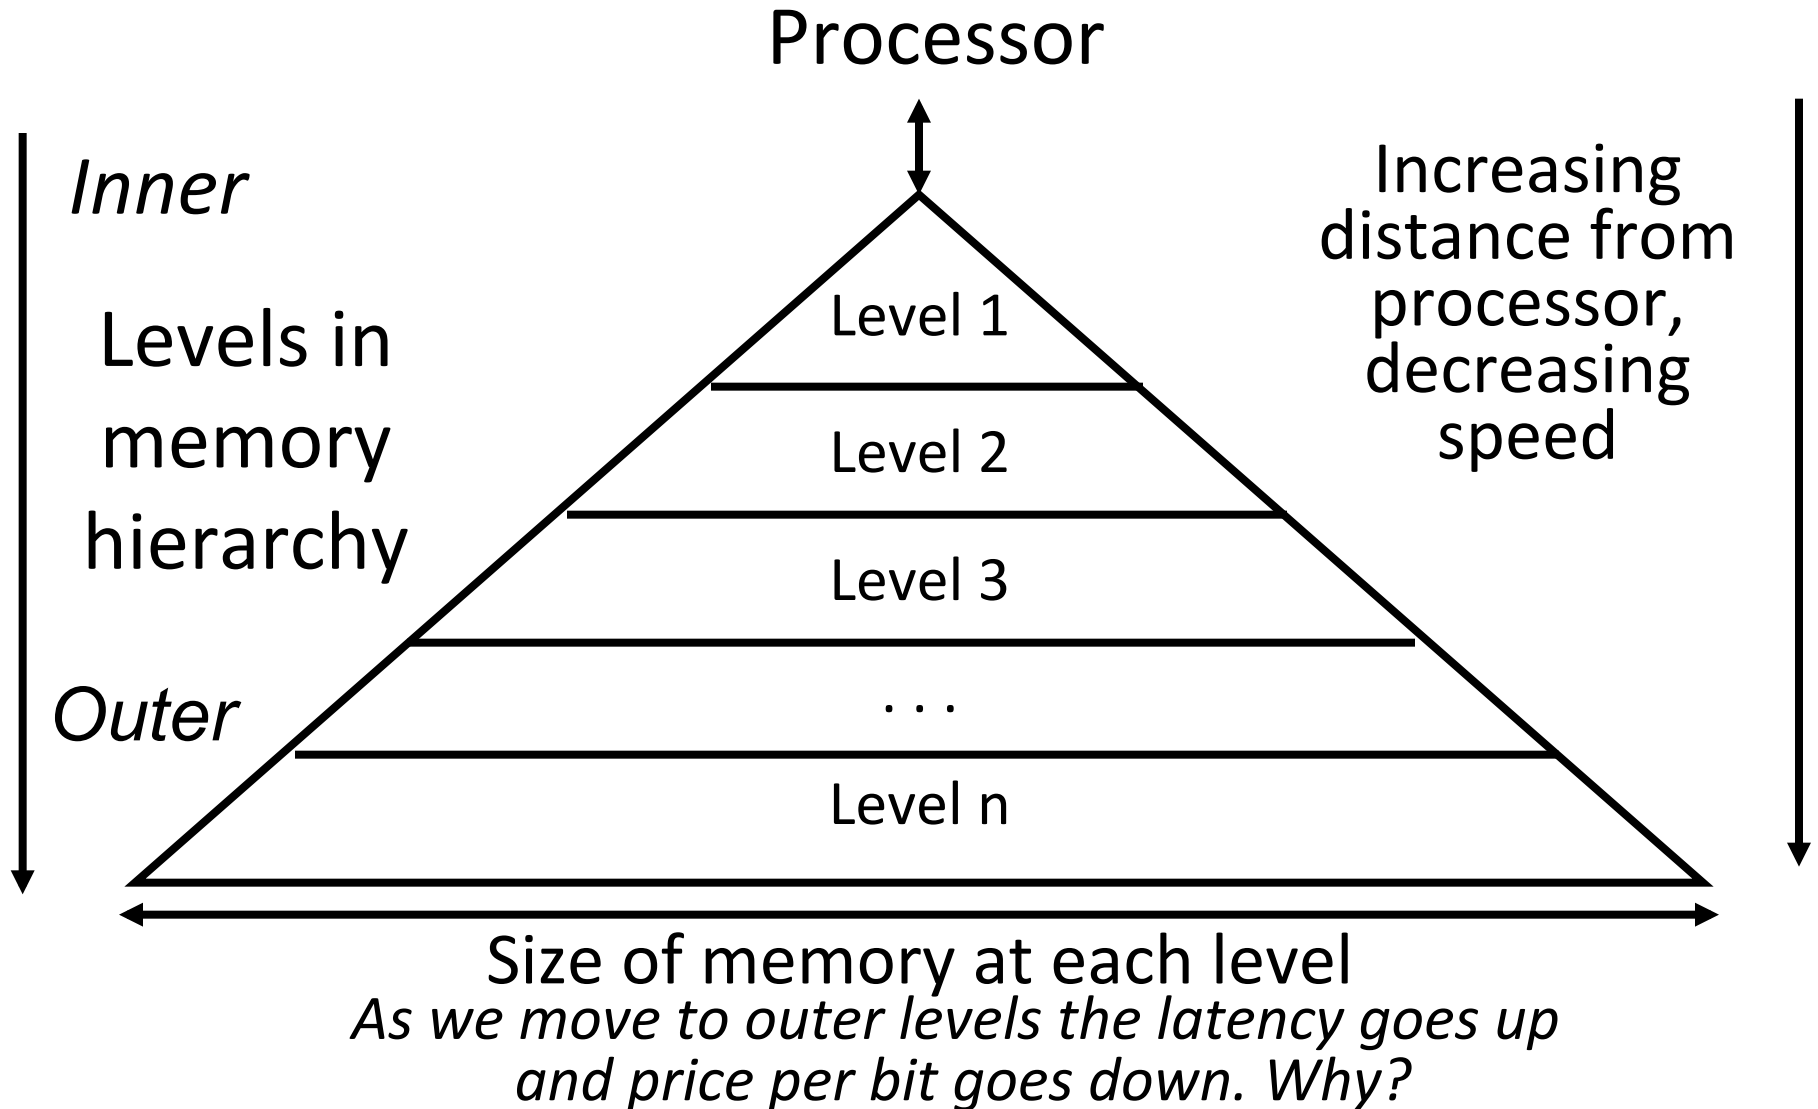
# Processor-DRAM Gap (latency)



1980 microprocessor executes ~one instruction in same time as DRAM access  
2015 microprocessor executes ~1000 instructions in same time as DRAM access



# Big Idea: Memory Hierarchy



# Library Analogy

- Writing a report based on books on reserve
  - E.g., works of J.D. Salinger
- Go to library to get reserved book and place on desk in library
- If need more, check them out and keep on desk
  - But don't return earlier books since might need them
- You hope this collection of ~10 books on desk enough to write report, despite 10 being only 0.00001% of books in UC Berkeley libraries

# Administrivia

- HW3 Out
- Proj2-2 Out
- Midterm regrades due today @ 23:59:59
- Register your Project 3 Teams

# In the News: RowHammer Exploit

## **Flipping Bits in Memory Without Accessing Them: An Experimental Study of DRAM Disturbance Errors**


Yoongu Kim<sup>1</sup> Ross Daly\* Jeremie Kim<sup>1</sup> Chris Fallin\* Ji Hye Lee<sup>1</sup>  
Donghyuk Lee<sup>1</sup> Chris Wilkerson<sup>2</sup> Konrad Lai Onur Mutlu<sup>1</sup>

<sup>1</sup>Carnegie Mellon University <sup>2</sup>Intel Labs

- CMU + Intel researchers found commercial DRAM chips susceptible to neighboring bits flipping if one row of memory accessed frequently
- Google Engineers figured out how to use this to gain root access on a machine! Almost all laptops susceptible, but server ECC memory helps reduce impact.

# Break

# Real Memory Reference Patterns

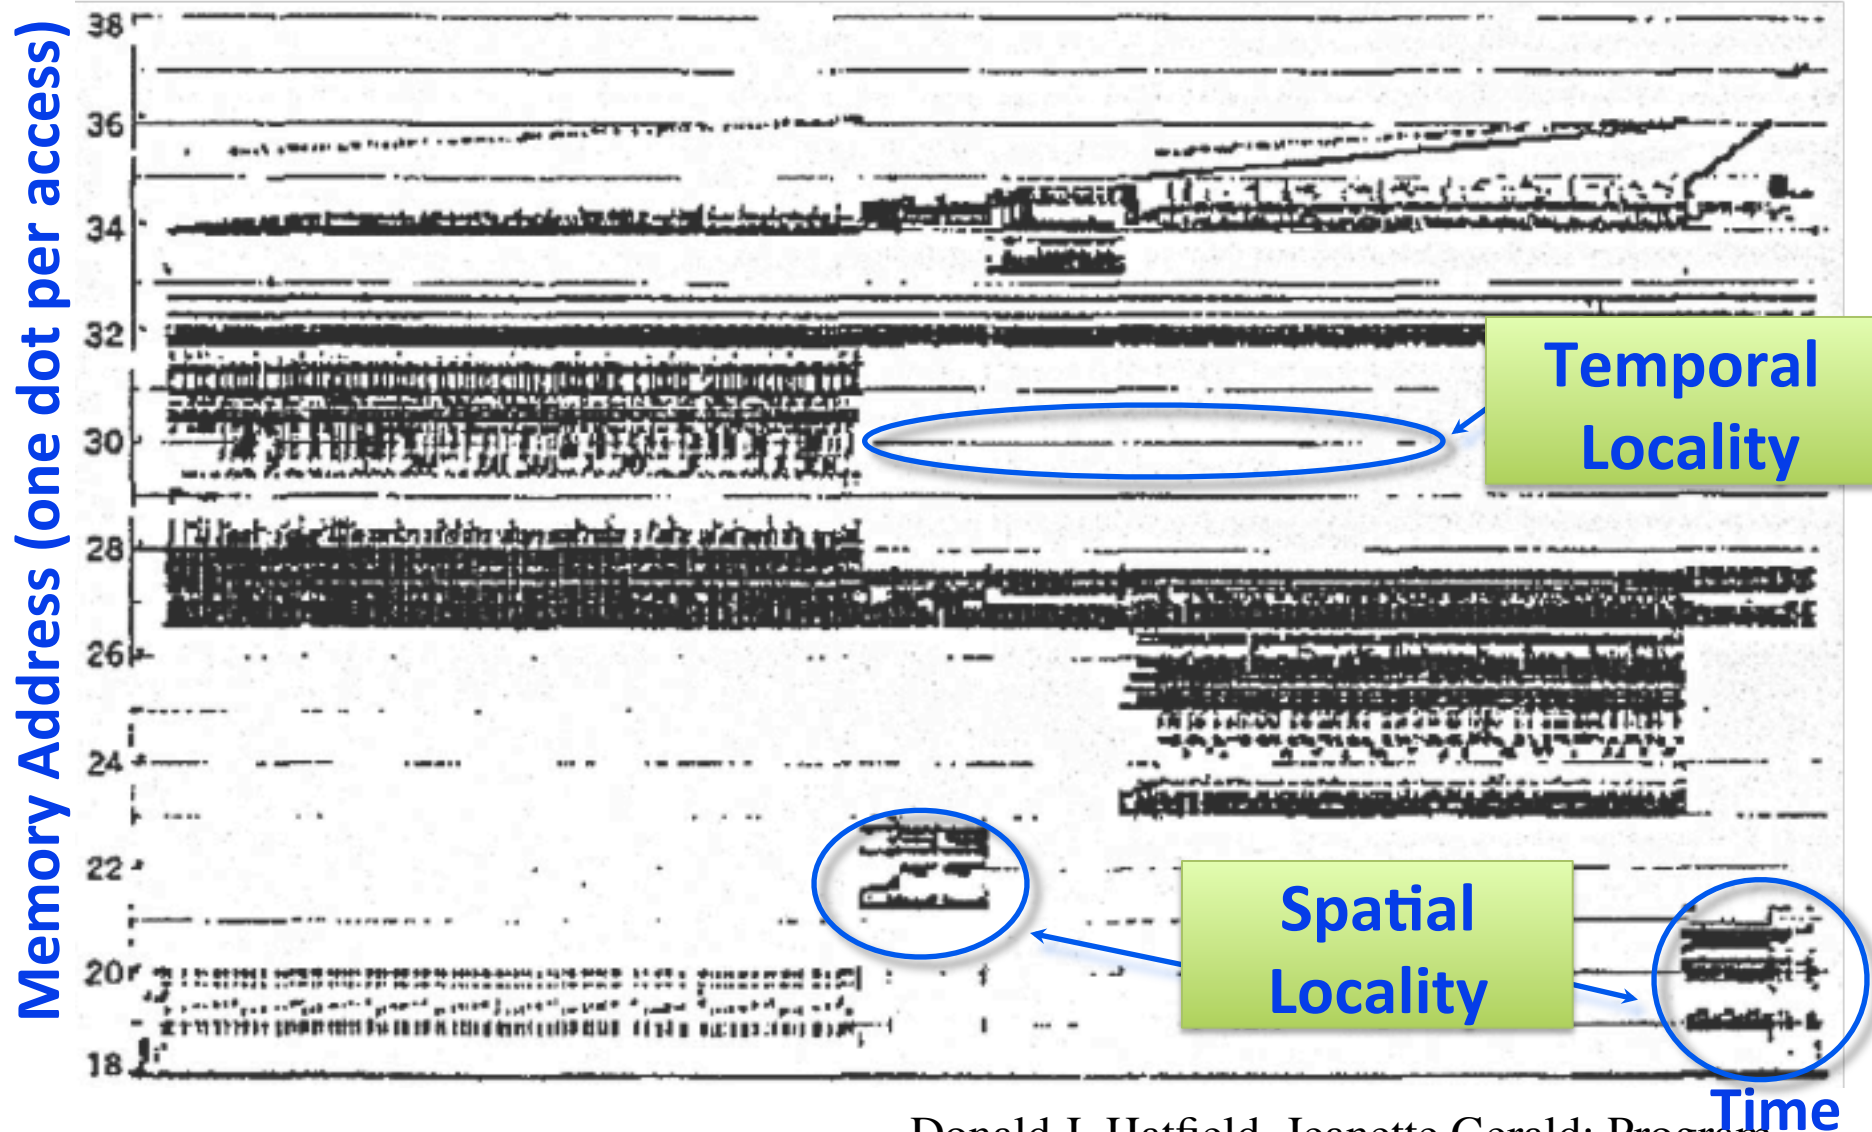


Donald J. Hatfield, Jeanette Gerald: Program Restructuring for Virtual Memory. IBM Systems Journal 10(3): 168-192 (1971)

# Big Idea: Locality

- *Temporal Locality* (locality in time)
  - Go back to same book on desktop multiple times
  - If a memory location is referenced, then it will tend to be referenced again soon
- *Spatial Locality* (locality in space)
  - When go to book shelf, pick up multiple books on J.D. Salinger since library stores related books together
  - If a memory location is referenced, the locations with nearby addresses will tend to be referenced soon

# Memory Reference Patterns



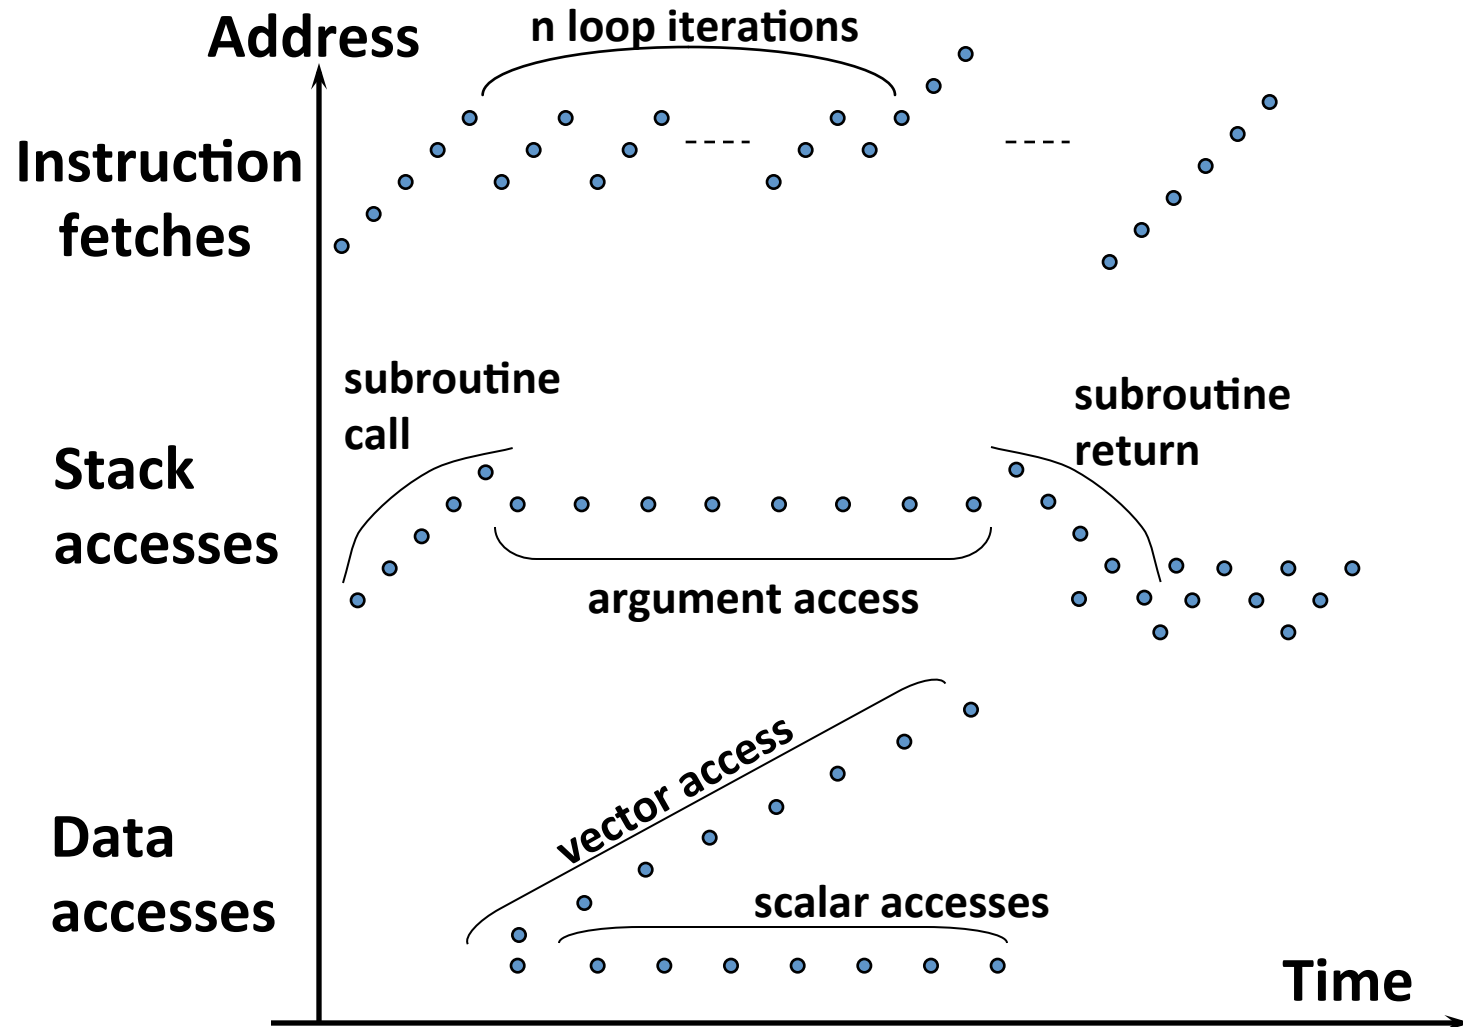
Donald J. Hatfield, Jeanette Gerald: Program Restructuring for Virtual Memory. IBM Systems Journal 10(3): 168-192 (1971)



# Principle of Locality

- *Principle of Locality*: Programs access small portion of address space at any instant of time
- What program structures lead to **temporal** and **spatial locality** in instruction accesses?
- In data accesses?

# Memory Reference Patterns



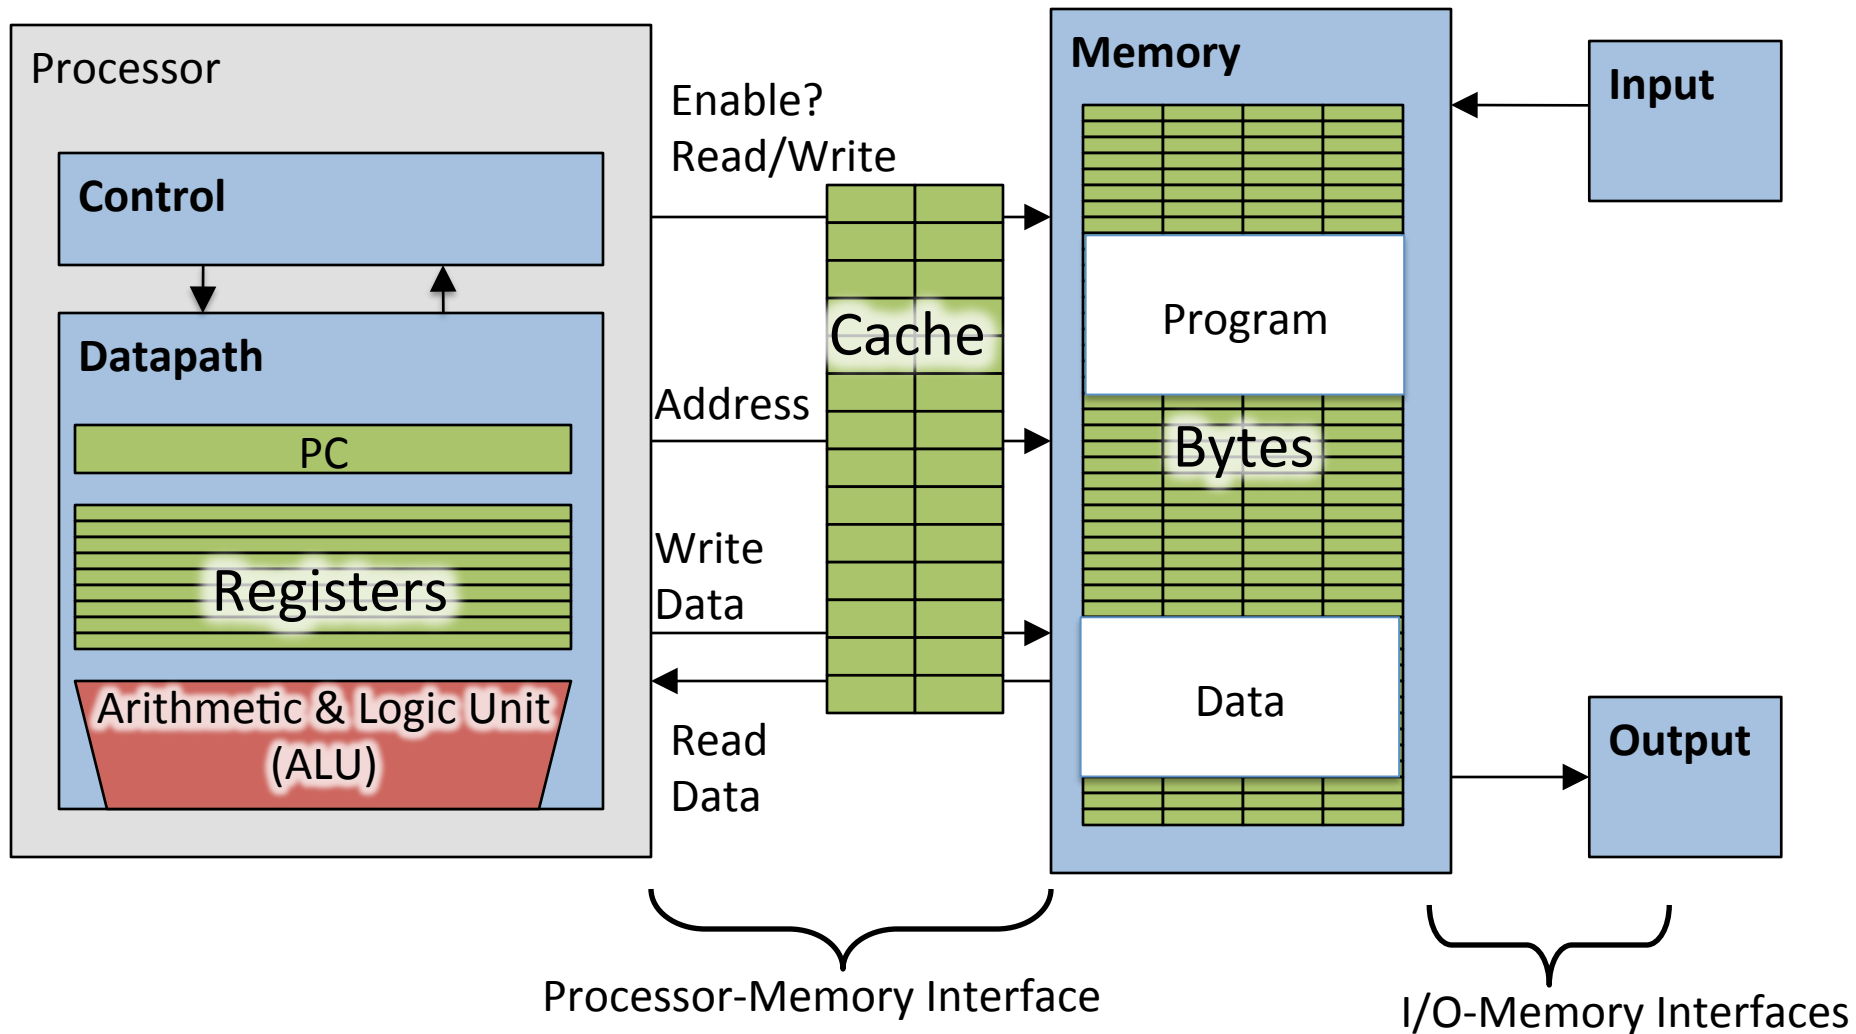
# Cache Philosophy

- Programmer-invisible hardware mechanism to give illusion of speed of fastest memory with size of largest memory
  - Works fine even if programmer has no idea what a cache is
  - However, performance-oriented programmers today sometimes “reverse engineer” cache design to design data structures to match cache
  - We’ll do this in Project 4

# Memory Access without Cache

- Load word instruction: `lw $t0, 0($t1)`
- `$t1` contains  $1022_{\text{ten}}$ , `Memory[1022] = 99`
  1. Processor issues address  $1022_{\text{ten}}$  to Memory
  2. Memory reads word at address  $1022_{\text{ten}}$  (99)
  3. Memory sends 99 to Processor
  4. Processor loads 99 into register `$t0`

# Adding Cache to Computer



# Memory Access with Cache

- Load word instruction: `lw $t0, 0($t1)`
- `$t1` contains  $1022_{ten}$ , `Memory[1022] = 99`
- With cache (similar to a hash)
  1. Processor issues address  $1022_{ten}$  to Cache
  2. Cache checks to see if has copy of data at address  $1022_{ten}$ 
    - 2a. If finds a match (Hit): cache reads 99, sends to processor
    - 2b. No match (Miss): cache sends address 1022 to Memory
      - I. Memory reads 99 at address  $1022_{ten}$
      - II. Memory sends 99 to Cache
      - III. Cache replaces word with new 99
      - IV. Cache sends 99 to processor
  3. Processor loads 99 into register `$t0`

# Cache “Tags”


- Need way to tell if have copy of location in memory so that can decide on hit or miss
- On cache miss, put memory address of block in “tag address” of cache block
  - 1022 placed in tag next to data from memory (99)

Tag	Data
252	12
1022	99
131	7
2041	20

From earlier instructions

# Anatomy of a 16 Byte Cache, 4 Byte Block

- Operations:
  1. Cache Hit
  2. Cache Miss
  3. Refill cache from memory
- Cache needs Address Tags to decide if Processor Address is a Cache Hit or Cache Miss
  - Compares all 4 tags





# Cache Replacement

- Suppose processor now requests location 511, which contains 11?
- Doesn't match any cache block, so must "evict" one resident block to make room
  - Which block to evict?
- Replace "victim" with new memory block at address 511

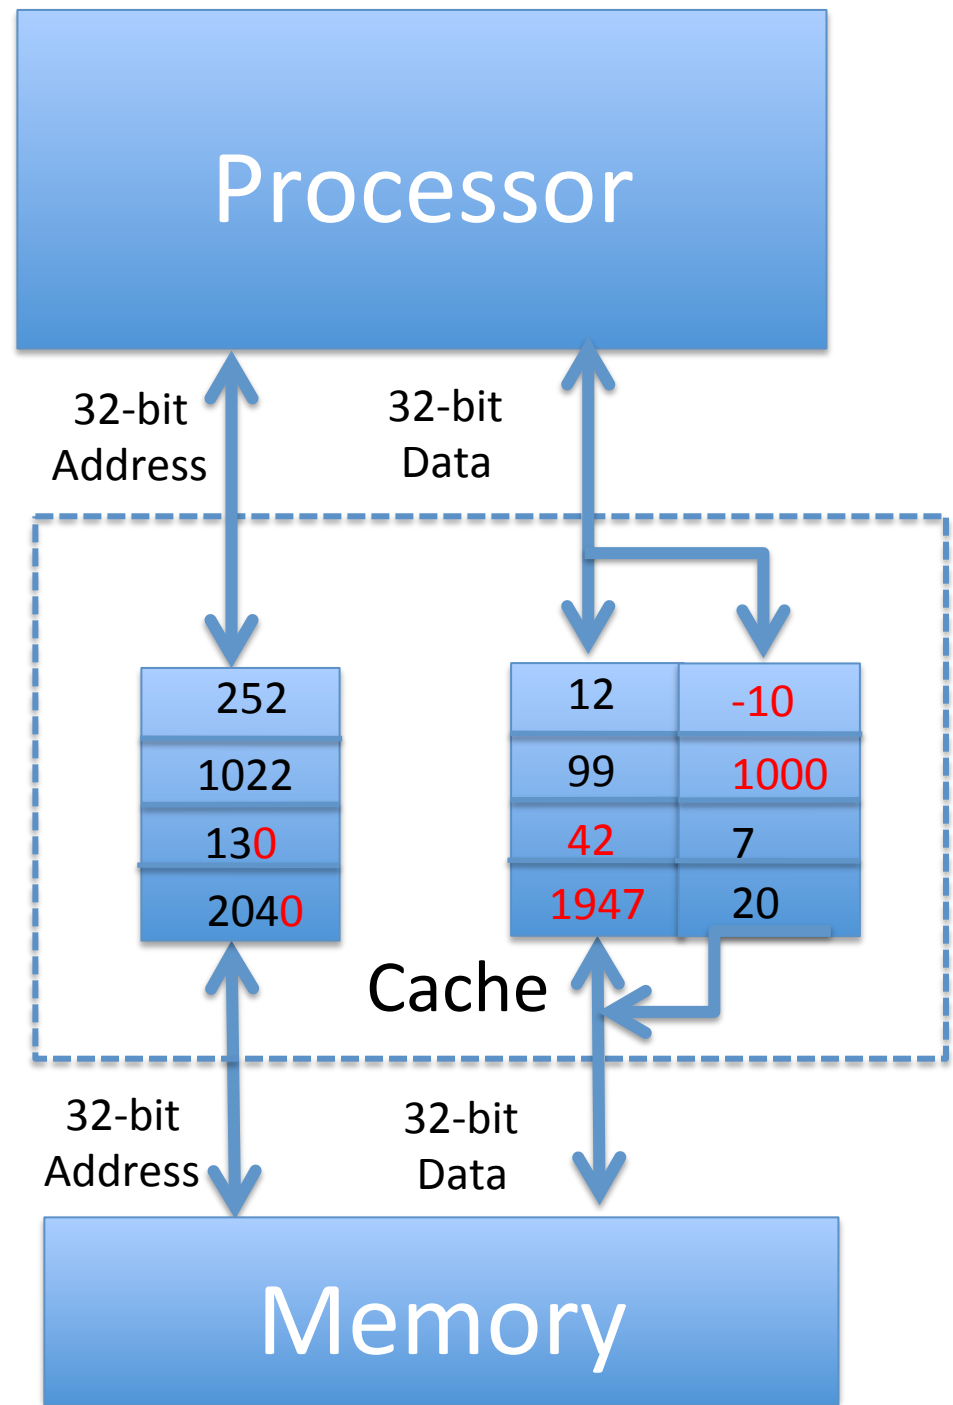
Tag	Data
252	12
1022	99
511	11
2041	20

# Block Must be Aligned in Memory

- Word blocks are aligned, so binary address of all words in cache always ends in  $00_{\text{two}}$
  - How to take advantage of this to save hardware and energy?
  - Don't need to compare last 2 bits of 32-bit byte address (comparator can be narrower)
- => Don't need to store last 2 bits of 32-bit byte address in Cache Tag (Tag can be narrower)

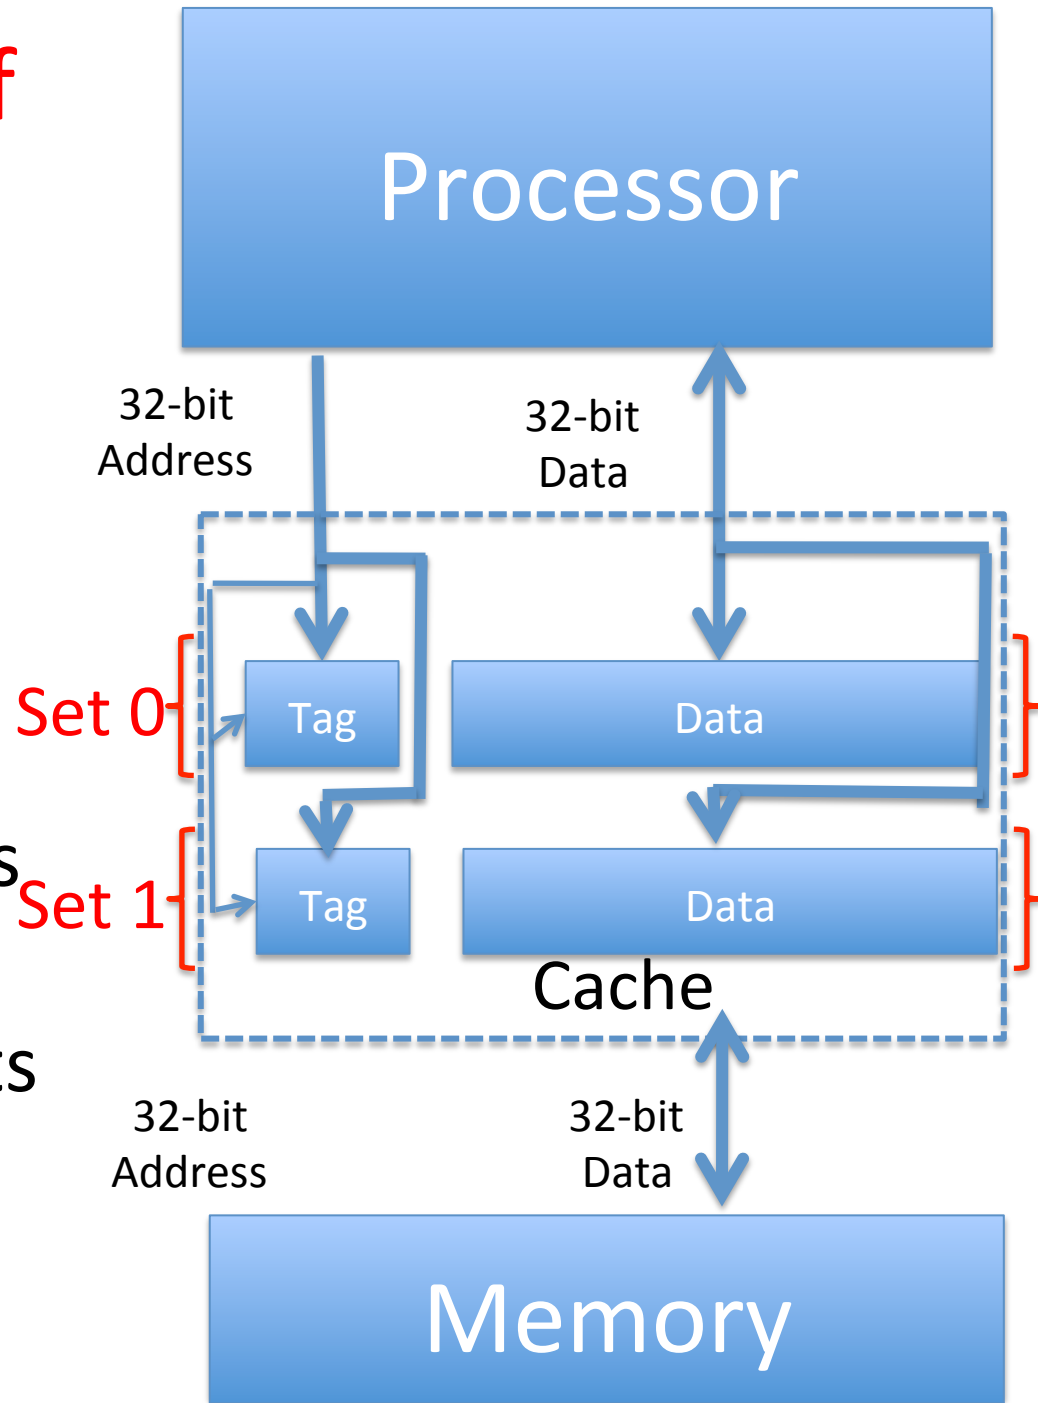
# Anatomy of a 32B Cache, 8B Block

- Blocks must be aligned in pairs, otherwise could get same word twice in cache
- ⇒ Tags only have even-numbered words
- ⇒ Last 3 bits of address always 000<sub>two</sub>
- ⇒ Tags, comparators can be narrower
- Can get hit for either word in block



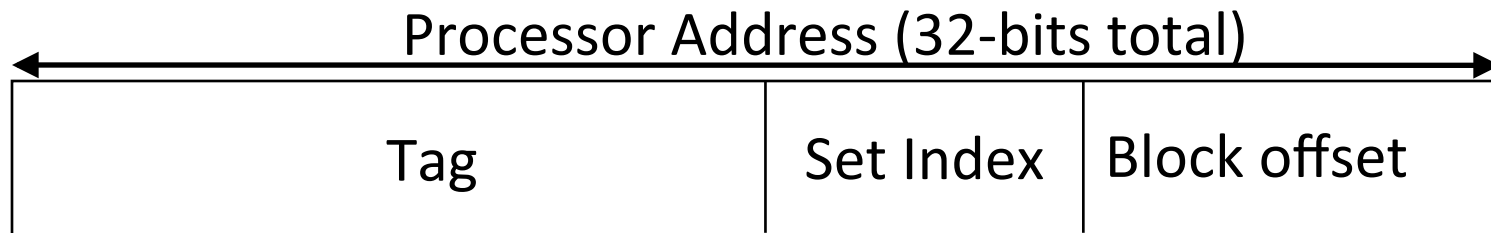
# Hardware Cost of Cache

- Need to compare every tag to the Processor address
- Comparators are expensive
- Optimization: 2 sets  $\Rightarrow \frac{1}{2}$  comparators
- 1 Address bit selects which set



# Processor Address Fields used by Cache Controller

- **Block Offset**: Byte address within block
- **Set Index**: Selects which set
- **Tag**: Remaining portion of processor address



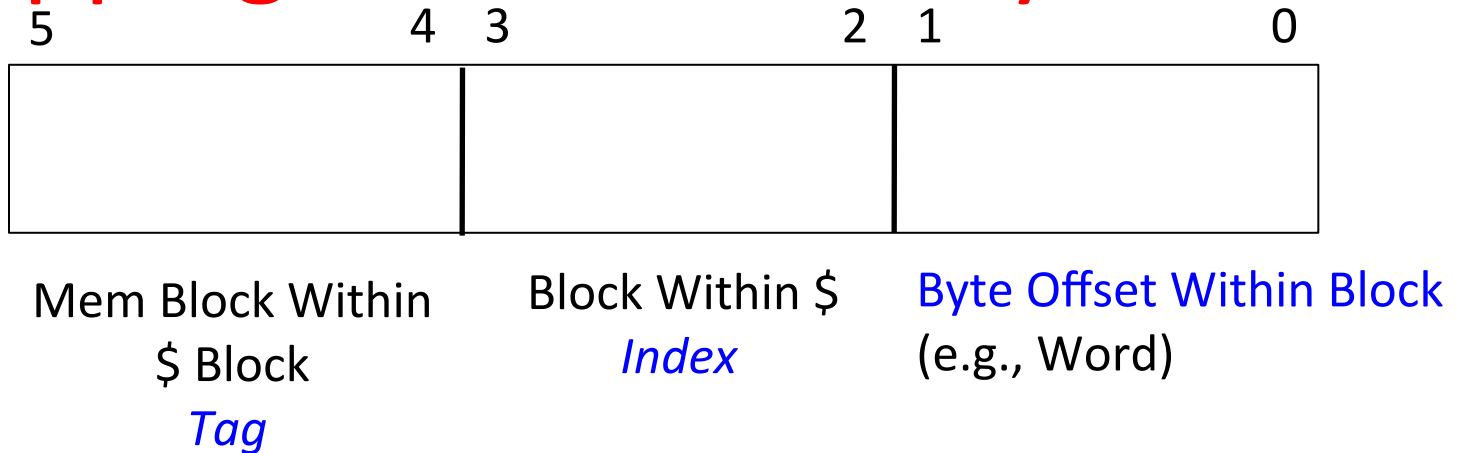
- Size of Index =  $\log_2$  (number of sets)
- Size of Tag = Address size – Size of Index –  $\log_2$  (number of bytes/block)

# What is limit to number of sets?

- Can save more comparators if have more than 2 sets
- Limit: As Many Sets as Cache Blocks – only needs one comparator!
- Called “Direct-Mapped” Design

Tag	Index	Block offset
-----	-------	--------------

# Mapping a 6-bit Memory Address



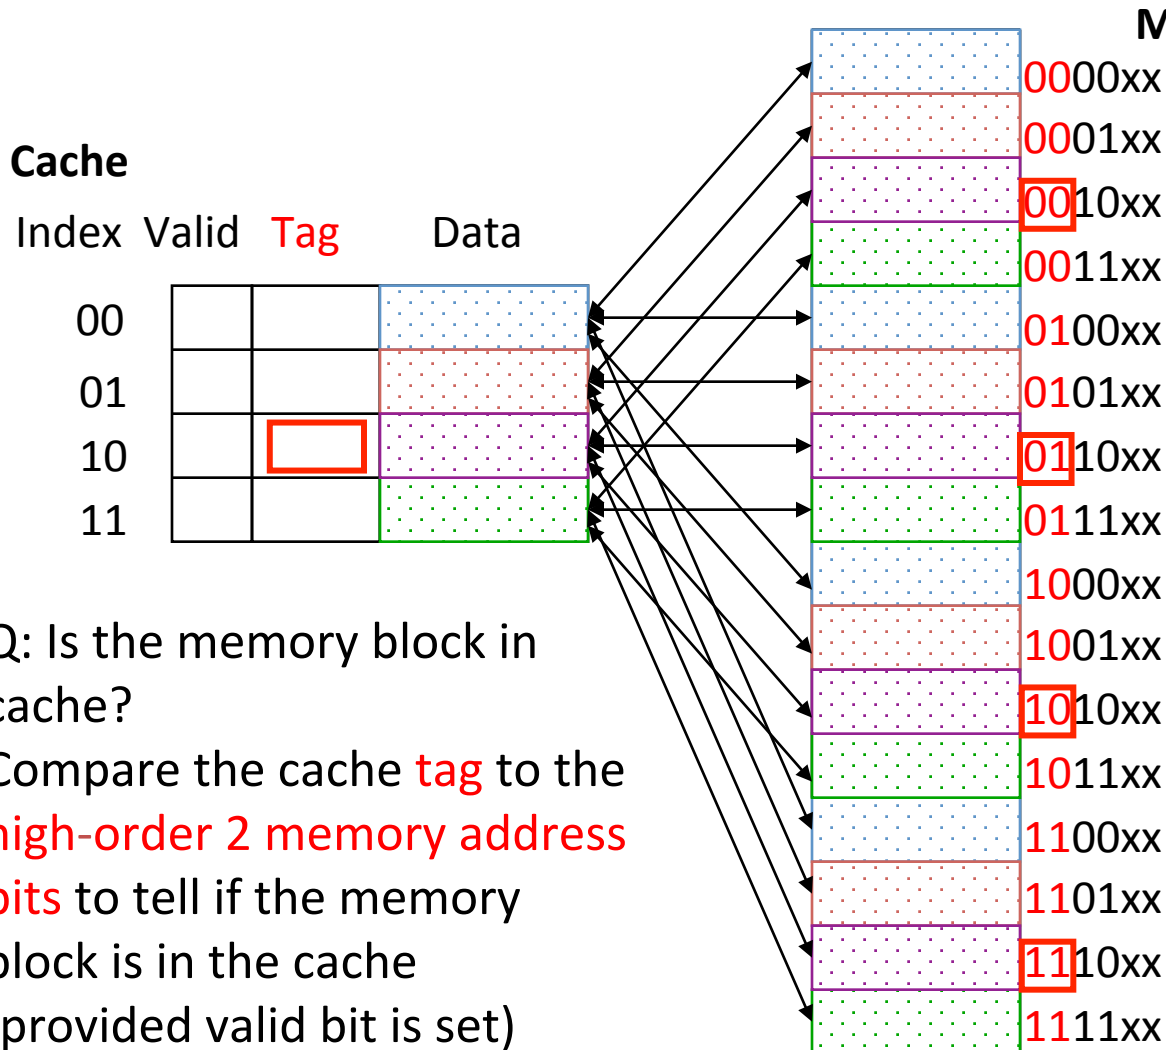
- In example, block size is 4 bytes/1 word (it could be multi-word)
- Memory and cache blocks are the same size, unit of transfer between memory and cache
- # Memory blocks >> # Cache blocks
  - 16 Memory blocks/16 words/64 bytes/6 bits to address all bytes
  - 4 Cache blocks, 4 bytes (1 word) per block
  - 4 Memory blocks map to each cache block
- Byte within block: low order two bits, ignore! (nothing smaller than a block)
- Memory block to cache block, aka *index*: middle two bits
- Which memory block is in a given cache block, aka *tag*: top two bits

# One More Detail: Valid Bit

- When start a new program, cache does not have valid information for this program
- Need an indicator whether this tag entry is valid for this program
- Add a “valid bit” to the cache tag entry
  - 0 => cache miss, even if by chance, address = tag
  - 1 => cache hit, if processor address = tag



# Caching: A Simple First Example



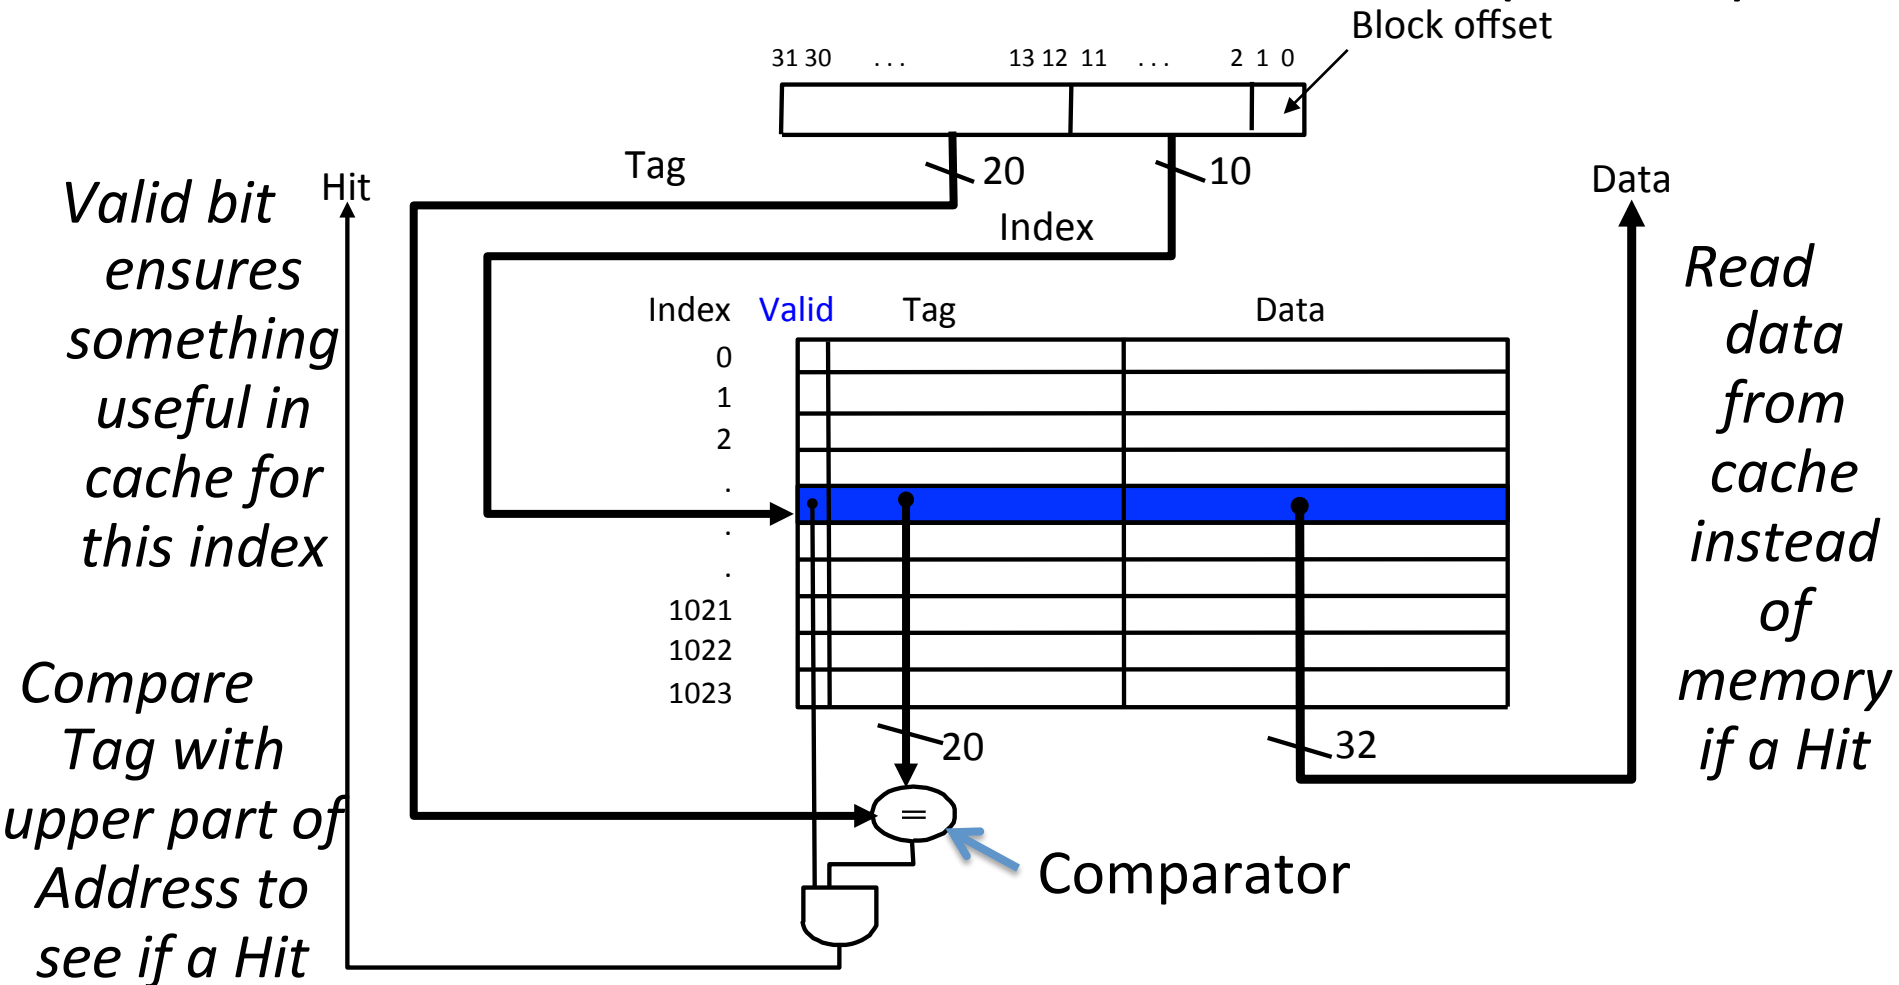
One word blocks  
Two low order bits (xx)  
define the byte in the  
block (32b words)

Q: Where in the cache is  
the mem block?

Use next 2 low-order  
memory address bits –  
the index – to determine  
which cache block (i.e.,  
modulo the number of  
blocks in the cache)

# Direct-Mapped Cache Example

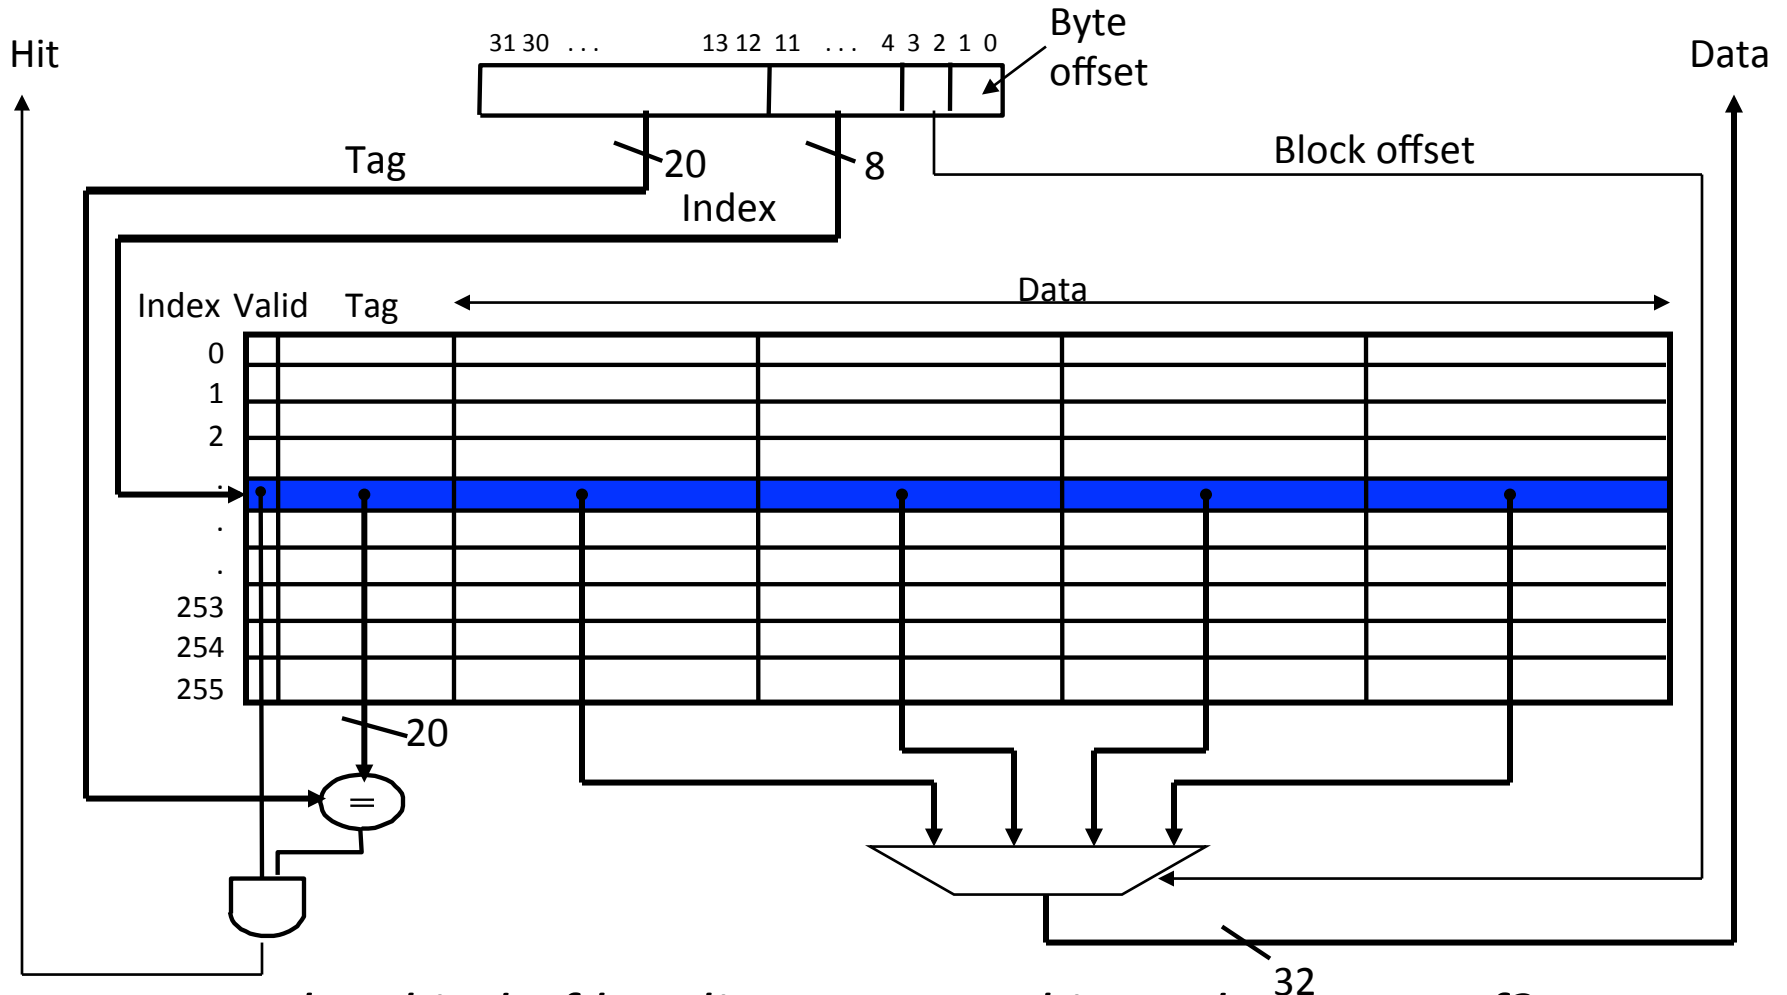
- One word blocks, cache size = 1K words (or 4KB)



*What kind of locality are we taking advantage of?*

# Multiword-Block Direct-Mapped Cache

- Four words/block, cache size = 1K words



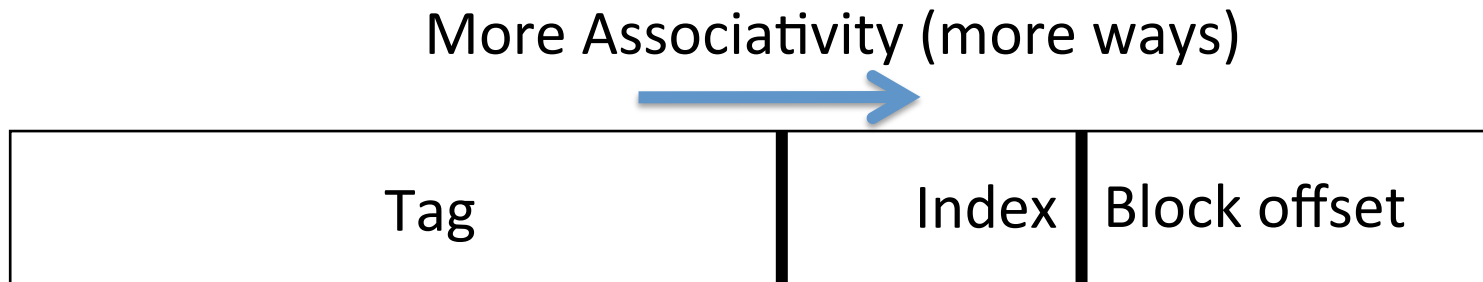
*What kind of locality are we taking advantage of?*

# Cache Names for Each Organization

- “Fully Associative”: Block can go anywhere
  - First design in lecture
  - Note: No Index field, but 1 comparator/block
- “Direct Mapped”: Block goes one place
  - Note: Only 1 comparator
  - Number of sets = number blocks
- “N-way Set Associative”: N places for a block
  - Number of sets = number of blocks / N
  - Fully Associative: N = number of blocks
  - Direct Mapped: N = 1

# Range of Set-Associative Caches

- For a fixed-size cache, each increase by a factor of 2 in associativity doubles the number of blocks per set (i.e., the number of “ways”) and halves the number of sets –
  - decreases the size of the index by 1 bit and increases the size of the tag by 1 bit



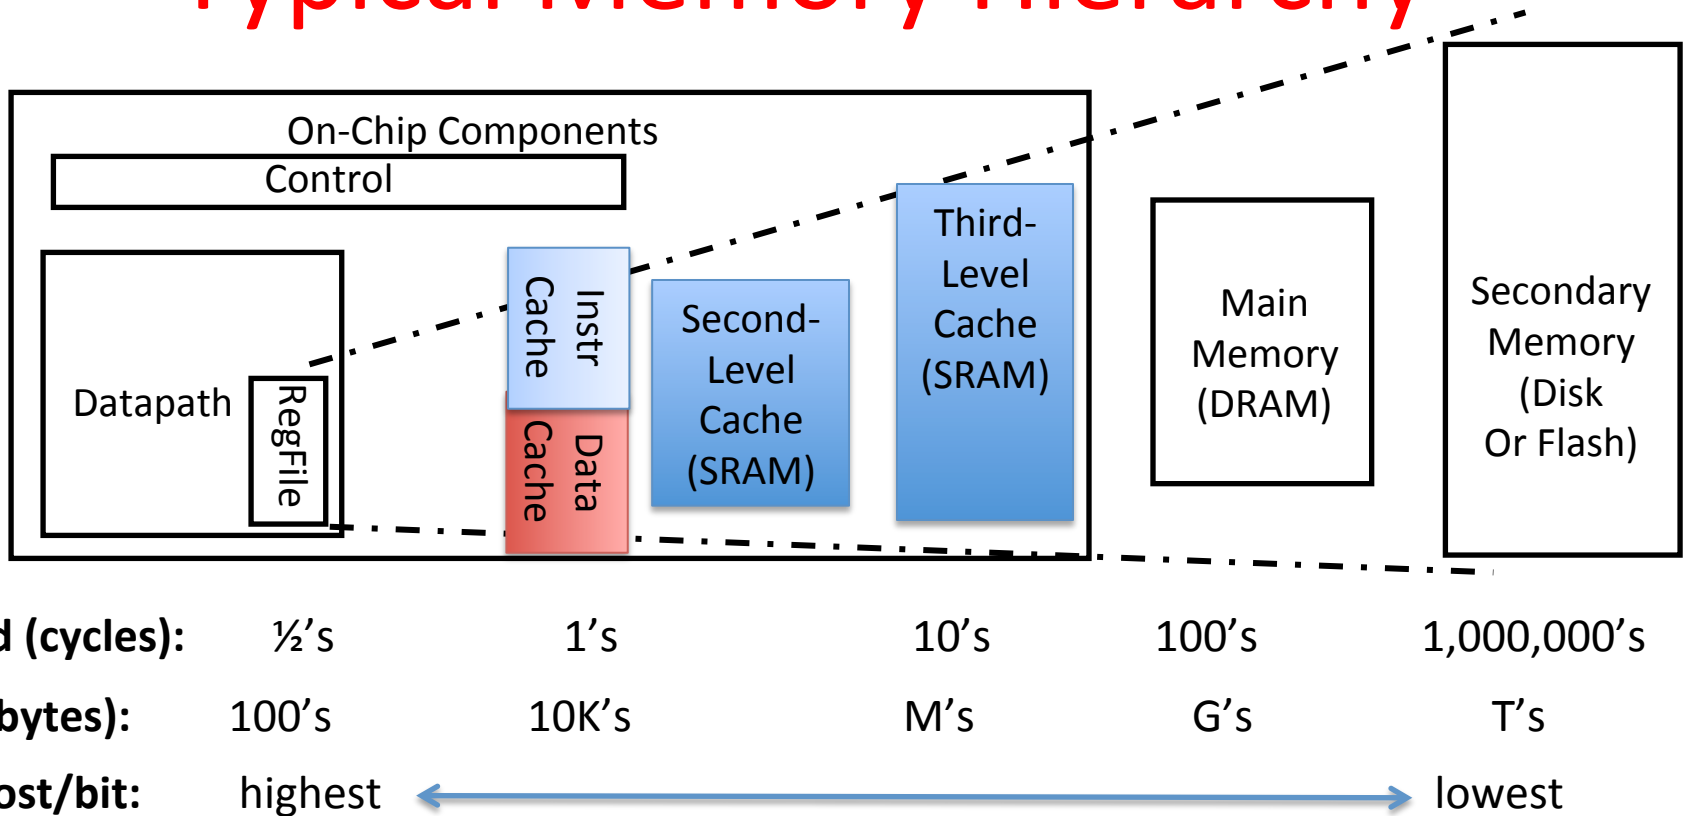
*Note: IBM persists in calling sets “ways” and ways “sets”. They’re wrong.*

# Clickers/Peer Instruction

- For a cache with constant total capacity, if we increase the number of ways by a factor of 2, which statement is false:
- A: The number of sets could be doubled
- B: The tag width could decrease
- C: The number of tags could stay the same
- D: The block size could be halved
- E: Tag width must increase

# Break

# Typical Memory Hierarchy



- **Principle of locality + memory hierarchy** presents programmer with  $\approx$  as much memory as is available in the *cheapest* technology at the  $\approx$  speed offered by the *fastest* technology

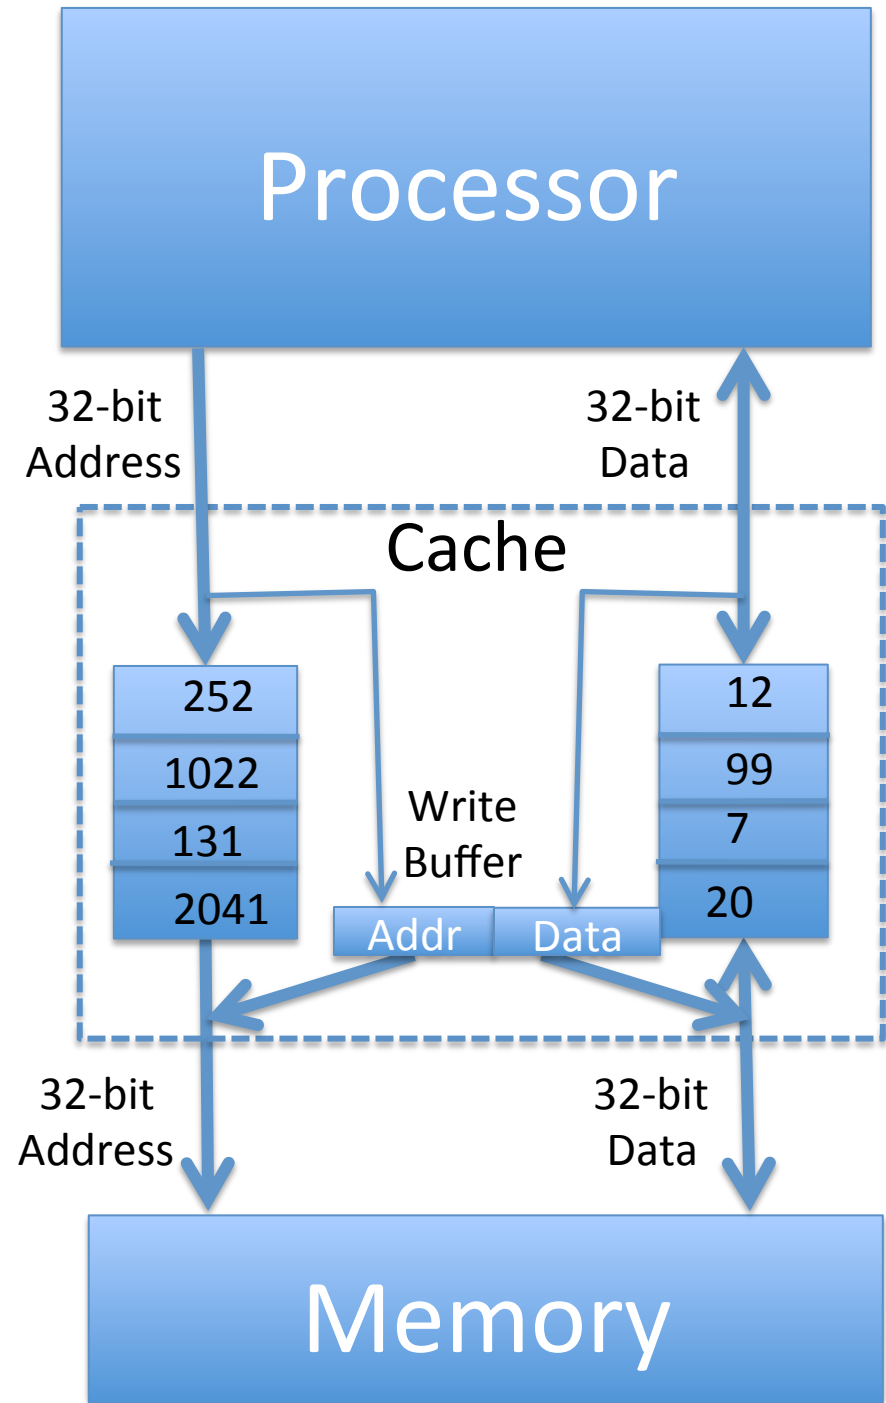


# Handling Stores with Write-Through

- Store instructions write to memory, changing values
  - Need to make sure cache and memory have same values on writes: 2 policies
- 1) **Write-Through Policy**: write cache and write *through* the cache to memory
- Every write eventually gets to memory
  - Too slow, so include Write Buffer to allow processor to continue once data in Buffer
  - Buffer updates memory in parallel to processor

# Write-Through Cache

- Write both values in cache and in memory
- Write buffer stops CPU from stalling if memory cannot keep up
- Write buffer may have multiple entries to absorb bursts of writes
- What if store misses in cache?

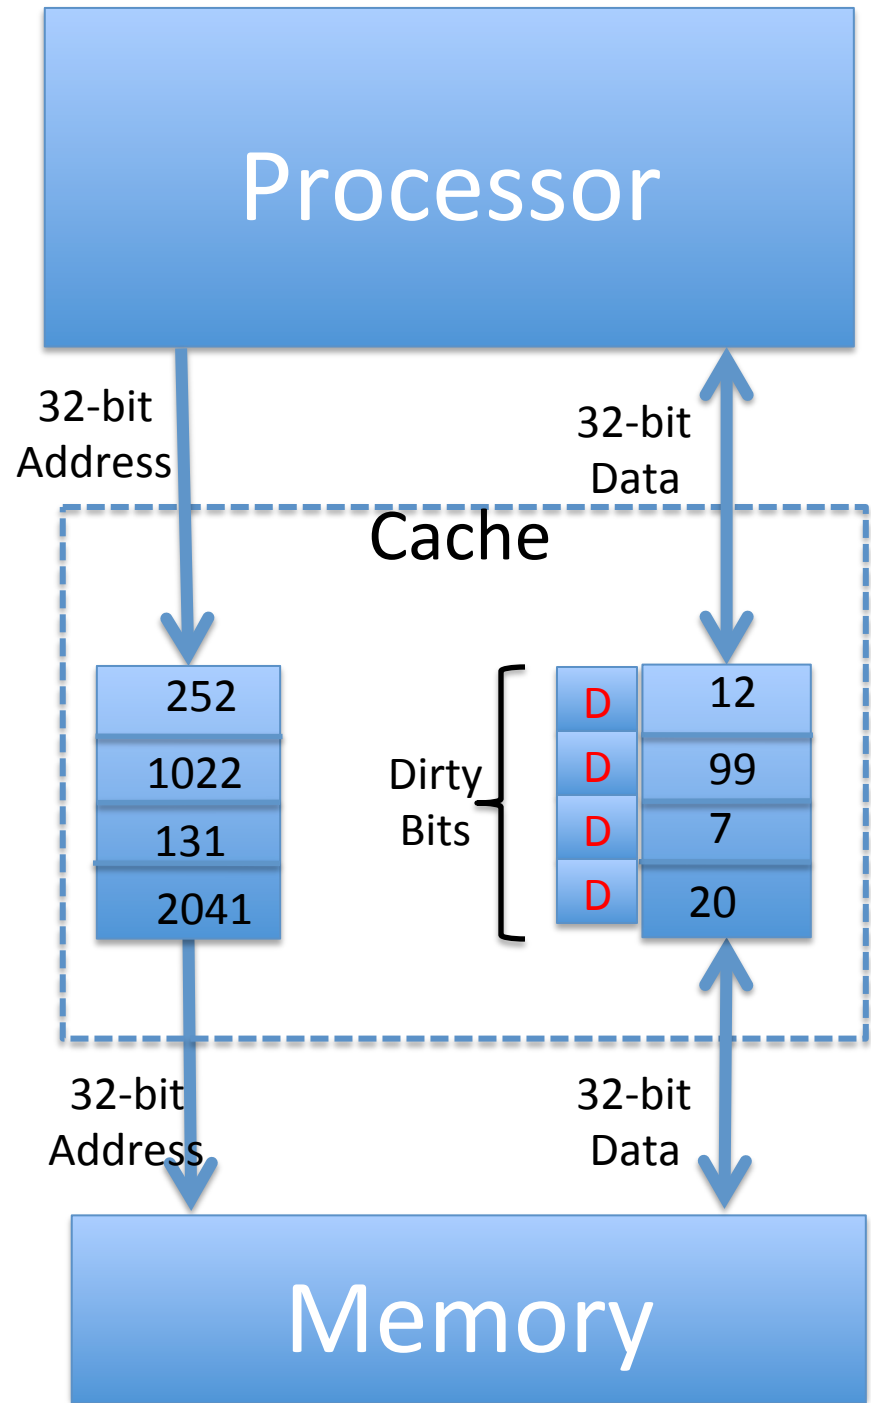


# Handling Stores with Write-Back

- 2) **Write-Back Policy**: write only to cache and then write cache block *back* to memory when evict block from cache
- Writes collected in cache, only single write to memory per block
  - Include bit to see if wrote to block or not, and then only write back if bit is set
    - Called “**Dirty**” bit (writing makes it “dirty”)

# Write-Back Cache

- Store/cache hit, write data in cache *only* & set dirty bit
  - Memory has stale value
- Store/cache miss, read data from memory, then update and set dirty bit
  - “Write-allocate” policy
- Load/cache hit, use value from cache
- On any miss, write back evicted block, only if dirty. Update cache with new block and clear dirty bit.



# Write-Through vs. Write-Back

- Write-Through:
  - Simpler control logic
  - More predictable timing  
simplifies processor control logic
  - Easier to make reliable, since memory always has copy of data
- Write-Back
  - More complex control logic
  - More variable timing (0,1,2 memory accesses per cache access)
  - Usually reduces write traffic
  - Harder to make reliable, sometimes cache has only copy of data

## And In Conclusion, ...

- Principle of Locality for Libraries /Computer Memory
- Hierarchy of Memories (speed/size/cost per bit) to Exploit Locality
- Cache – copy of data lower level in memory hierarchy
- Direct Mapped to find block in cache using Tag field and Valid bit for Hit
- Cache design choice:
  - Write-Through vs. Write-Back