Conditional Joints

Conditional Distributions

Before we looked at conditional probabilities for events. Here we formally go over conditional probabilities for random variables. The equations for both the discrete and continuous case are intuitive extensions of our understanding of conditional probability:

Discrete

The conditional probability mass function (PMF) for the discrete case:

$$p_{X|Y}(x|y) = P(X = x|Y = y) = \frac{P(X = x, Y = y)}{P(Y = y)} = \frac{P_{X,Y}(x,y)}{p_Y(y)}$$

The conditional cumulative density function (CDF) for the discrete case:

$$F_{X|Y}(a|y) = P(X \le a|Y = y) = \frac{\sum_{x \le a} p_{X,Y}(x,y)}{p_Y(y)} = \sum_{x \le a} p_{X|Y}(x|y)$$

Continuous

The conditional probability density function (PDF) for the continuous case:

$$f_{X|Y}(x|y) = \frac{f_{X,Y}(x,y)}{f_Y(y)}$$

The conditional cumulative density function (CDF) for the continuous case:

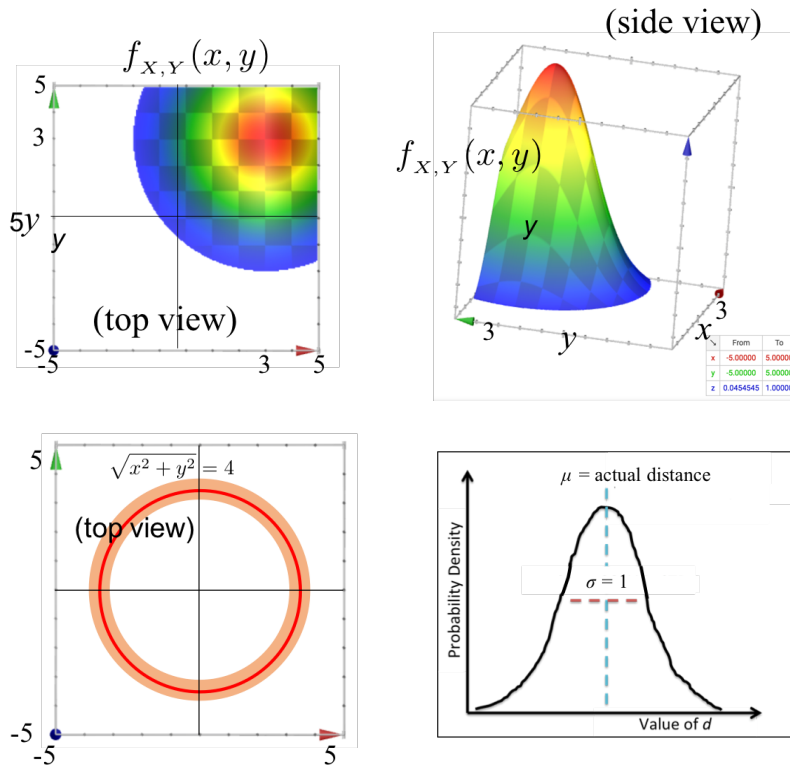$$F_{X|Y}(a|y) = P(X \le a|Y = y) = \int_{-\infty}^{a} f_{X|Y}(x|y)dx$$

Example

In this example we are going to explore the problem of tracking an object in 2D space. The object exists at some (x,y) location, however we are not sure exactly where! Thus we are going to use random variables X and Y to represent location.

We have a prior belief about where the object is. In this example our prior both X and Y as normals which are independently distributed with mean 3 and variance 2. First let's write the prior belief as a joint probability density function

$$
\begin{aligned}
f(X = x, Y = y) &= f(X = x) \cdot f(Y = y) &&\text{In the prior X and Y are independent} \\
&= \frac{1}{\sqrt{2 \cdot 4 \cdot \pi}} \cdot e^{-\frac{(x-3)^2}{2 \cdot 4}} \cdot \frac{1}{\sqrt{2 \cdot 4 \cdot \pi}} \cdot e^{-\frac{(y-3)^2}{2 \cdot 4}} &&\text{Using the PDF equation for normals} \\
&= K_1 \cdot e^{-\frac{(x-3)^2 + (y-3)^2}{8}} &&\text{All constants are put into } K_1
\end{aligned}
$$

This combinations of normals is called a bivariate distribution. Here is a visualization of the PDF of our prior.

The interesting part about tracking an object is the process of updating your belief about it's location based on an observation. Let's say that we get an instrument reading from a sonar that is sitting on the origin. The

$f_{X,Y}(x,y)$ (top view)

(side view)

$f_{X,Y}(x,y)$

	From	To
x	-5.00000	5.00000
y	-5.00000	5.00000
z	0.0454545	1.00000

$\sqrt{x^2+y^2}=4$ (top view)

μ = actual distance

$\sigma = 1$

Probability Density

Value of d

instrument reports that the object is 4 units away. Our instrument is not perfect: if the true distance was t units away, than the instrument will give a reading which is normally distributed with mean t and variance 1. Let's visualize the observation:

Based on this information about the noisiness of our prior, we can compute the conditional probability of seeing a particular distance reading D, given the true location of the object X, Y. If we knew the object was at location (x, y), we could calculate the true distance to the origin $\sqrt{x^2+y^2}$ which would give us the mean for the instrument Gaussian:

$$f(D = d|X = x, Y = y) = \frac{1}{\sqrt{2 \cdot 1 \cdot \pi}} \cdot e^{-\frac{\left(d-\sqrt{x^2+y^2}\right)^2}{2 \cdot 1}} \qquad \text{Normal PDF where } \mu = \sqrt{x^2+y^2}$$

$$= K_2 \cdot e^{-\frac{\left(d-\sqrt{x^2+y^2}\right)^2}{2 \cdot 1}} \qquad \text{All constants are put into } K_2$$

How about we try this out on actual numbers. How much more likely is an instrument reading of 1 compared to 2, given that the location of the object is at (1, 1)?

$$\frac{f(D = 1|X = 1, Y = 1)}{f(D = 2|X = 1, Y = 1)} = \frac{K_2 \cdot e^{-\frac{\left(1-\sqrt{1^2+1^2}\right)^2}{2 \cdot 1}}}{K_2 \cdot e^{-\frac{\left(2-\sqrt{1^2+1^2}\right)^2}{2 \cdot 1}}} \qquad \text{Substituting into the conditional PDF of D}$$

$$= \frac{e^0}{e^{-1/2}} \approx 1.65 \qquad \text{Notice how the } K_2 \text{ cancel out}$$

At this point we have a prior belief and we have an observation. We would like to compute an updated belief, given that observation. This is a classic Bayes' formula scenario. We are using joint continuous variables, but

2

that doesn't change the math much, it just means we will be dealing with densities instead of probabilities:

$$f(X = x, Y = y | D = 4) = \frac{f(D = 4 | X = x, Y = y) \cdot f(X = x, Y = y)}{f(D = 4)} \quad \text{Bayes using densities}$$

$$= \frac{K_1 \cdot e^{-\frac{[4 - \sqrt{x^2 + y^2})^2]}{2}} \cdot K_2 \cdot e^{-\frac{[(x-3)^2 + (y-3)^2]}{8}}}{f(D = 4)} \quad \text{Substituting for prior and update}$$

$$= \frac{K_1 \cdot K_2}{f(D = 4)} \cdot e^{-\left[\frac{[4 - \sqrt{x^2 + y^2})^2]}{2} + \frac{[(x-3)^2 + (y-3)^2]}{8}\right]} \quad f(D = 4) \text{ is a constant w.r.t. } (x, y)$$

$$= K_3 \cdot e^{-\left[\frac{(4 - \sqrt{x^2 + y^2})^2}{2} + \frac{[(x-3)^2 + (y-3)^2]}{8}\right]} \quad K_3 \text{ is a new constant}$$

Wow! That looks like a pretty interesting function! You have successfully computed the updated belief. Let's see what it looks like. Here is a figure with our prior on the left and the posterior on the right: How beautiful

Prior   Posterior
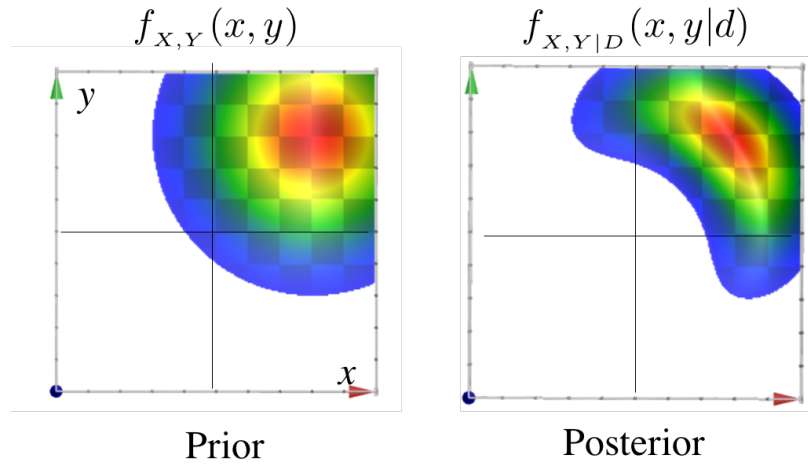
is that! Its like a 2D normal distribution merged with a circle. But wait, what about that constant! We do not know the value of K_3 and that is not a problem for two reasons: the first reason is that if we ever want to calculate a relative probability of two locations, K_3 will cancel out. The second reason is that if we really wanted to know what K_3 was, we could solve for it.

This math is used every day in millions of applications. If there are multiple observations the equations can get truly complex (even worse than this one). To represent these complex functions often use an algorithm called particle filtering.

Example

Let's say we have two independent random Poisson variables for requests received at a web server in a day: X = # requests from humans/day, $X \sim Poi(\lambda_1)$ and Y = # requests from bots/day, $Y \sim Poi(\lambda_2)$. Since the convolution of Poisson random variables is also a Poisson we know that the total number of requests $(X + Y)$ is also a Poisson $(X + Y) \sim Poi(\lambda_1 + \lambda_2)$. What is the probability of having k human requests on a particular

day given that there were n total requests?

$$P(X = k | X + Y = n) = \frac{P(X = k, Y = n - k)}{P(X + Y = n)} = \frac{P(X = k)P(Y = n - k)}{P(X + Y = n)}$$

$$= \frac{e^{-\lambda_1}\lambda_1^k}{k!} \cdot \frac{e^{-\lambda_2}\lambda_2^{n-k}}{(n-k)!} \cdot \frac{n!}{e^{1(\lambda_1 + \lambda_2)}(\lambda_1 + \lambda_2)^n}$$

$$= \binom{n}{k}\left(\frac{\lambda_1}{\lambda_1 + \lambda_2}\right)^k \left(\frac{\lambda_2}{\lambda_1 + \lambda_2}\right)^{n-k}$$

$$\sim Bin\left(n, \frac{\lambda_2}{\lambda_1 + \lambda_2}\right)$$