

Lecture 12

Bellman-Ford, Floyd-Warshall,
and Dynamic Programming!

Announcements

- HW6 out today!
- We are almost done grading the midterm – grades will be released soon.
 - Please follow standard procedure for regrade requests.
- I think the midterm was hard!
 - Great job!

Midterm Feedback

- I messed up.
 - Thank you to those who respectfully pointed out that there is actually some guidance from Stanford about timed take-home midterms.
- I think that students followed the honor code.
 - The grade distribution seems about right for a timed exam.
- However!
 - I don't want to go against Stanford's guidance on this, and I do want to address the legitimate concerns raised by students.
 - So...

New plan

- First, we will generate final letter grades as discussed on the website.
- Second, we will generate a *second* set of final letter grades as discussed on the website, except we will drop the midterm.
- You will receive the *maximum* of these two letter grades.

This is a Pareto-improving change! No one will receive a worse letter grade than they would under the original grading scheme!



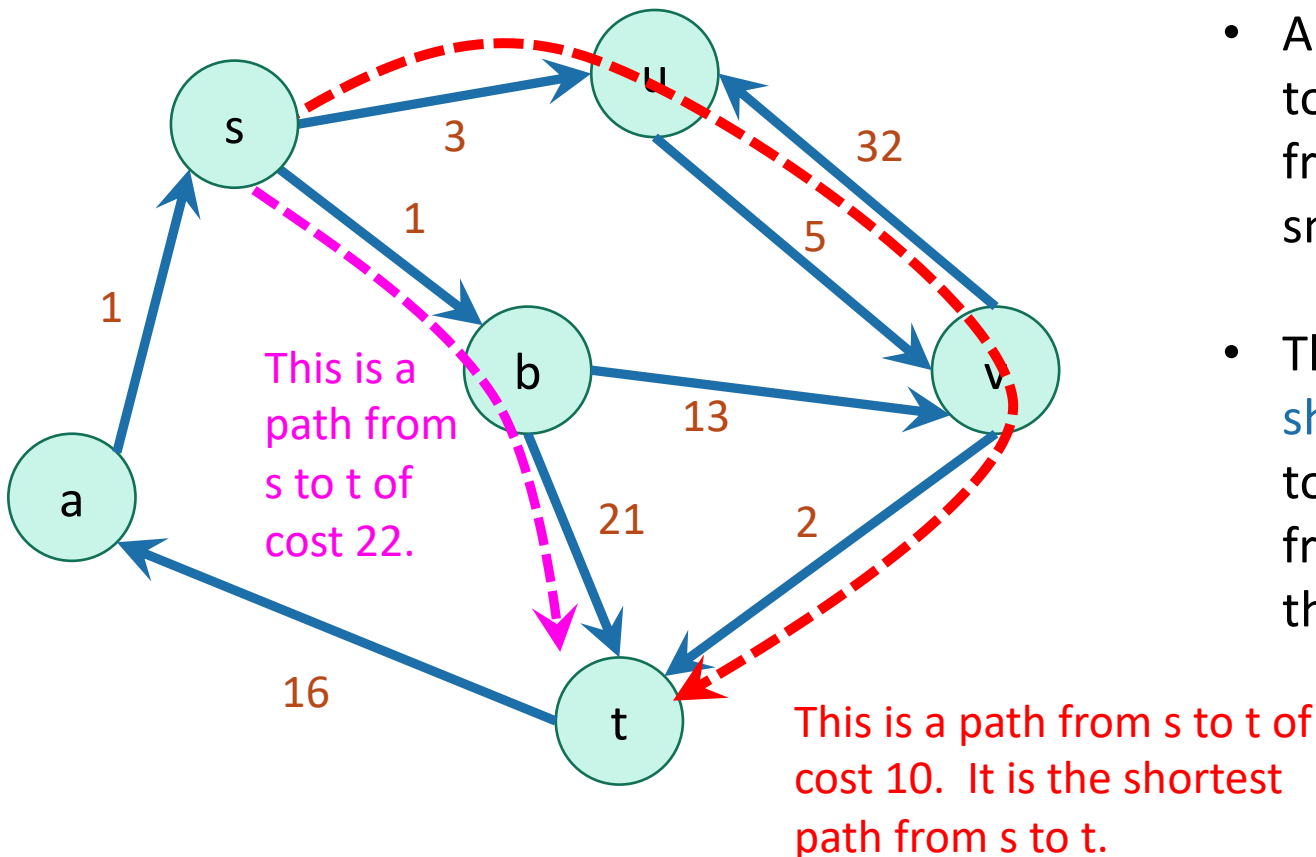
If you have questions, comments, or concerns about this policy, please post privately on Piazza or email the staff list.

Today

- Bellman-Ford Algorithm
- Bellman-Ford is a special case of ***Dynamic Programming!***
- What is dynamic programming?
 - Warm-up example: Fibonacci numbers
- Another example:
 - Floyd-Warshall Algorithm

Recall


- A weighted directed graph:



- Weights on edges represent **costs**.
- The **cost of a path** is the sum of the weights along that path.
- A **shortest path** from s to t is a directed path from s to t with the smallest cost.
- The **single-source shortest path problem** is to find the shortest path from s to v for all v in the graph.

Last time

- Dijkstra's algorithm!
 - Solves the single-source shortest path problem in weighted graphs.



Dijkstra Drawbacks


- Needs **non-negative edge weights**.
- If the weights change, we need to re-run the whole thing.

Bellman-Ford algorithm

- (-) Slower than Dijkstra's algorithm
- (+) Can handle negative edge weights.
 - Can be useful if you want to say that some edges are actively good to take, rather than costly.
 - Can be useful as a building block in other algorithms.
- (+) Allows for some flexibility if the weights change.
 - We'll see what this means later

Aside: Negative Cycles

- A **negative cycle** is a cycle whose edge weights sum to a negative number.
- Shortest paths aren't defined when there are negative cycles!



The shortest path from A to B has cost...negative infinity?


Bellman-Ford algorithm

- (-) Slower than Dijkstra's algorithm
- (+) Can handle negative edge weights.
 - Can **detect** negative cycles!
 - Can be useful if you want to say that some edges are actively good to take, rather than costly.
 - Can be useful as a building block in other algorithms.
- (+) Allows for some flexibility if the weights change.
 - We'll see what this means later

Bellman-Ford vs. Dijkstra

- Dijkstra:
 - Find the u with the smallest $d[u]$
 - Update u 's neighbors: $d[v] = \min(d[v], d[u] + w(u,v))$
- Bellman-Ford:
 - Don't bother finding the u with the smallest $d[u]$
 - Everyone updates!

Bellman-Ford

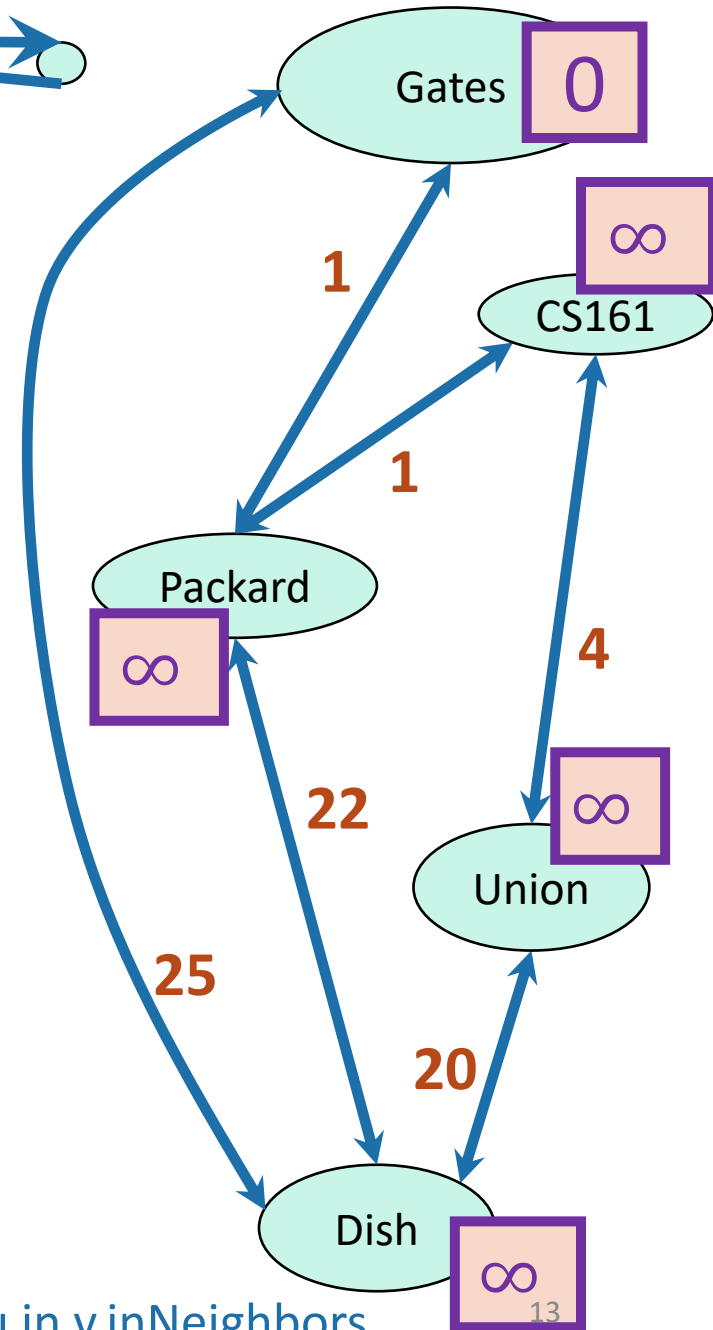


How far is a node from Gates?

	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$					
$d^{(2)}$					
$d^{(3)}$					
$d^{(4)}$					

- For $i=0, \dots, n-2$:
 - For v in V :
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v] , d^{(i)}[u] + w(u,v))$

where we are also taking the min over all u in $v.inNeighbors$

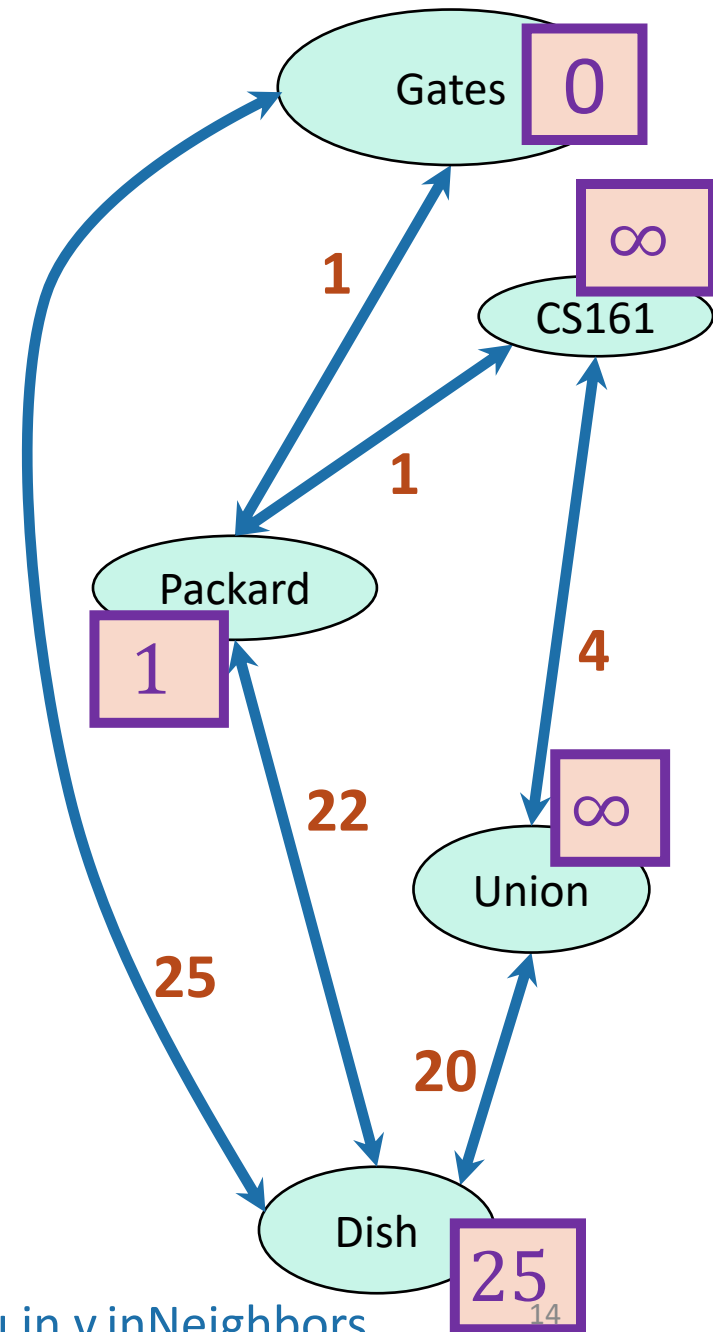


Bellman-Ford

How far is a node from Gates?

	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	25
$d^{(2)}$					
$d^{(3)}$					
$d^{(4)}$					

- For $i=0, \dots, n-2$:
 - For v in V :
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v] , d^{(i)}[u] + w(u,v))$
 where we are also taking the min over all u in $v.inNeighbors$

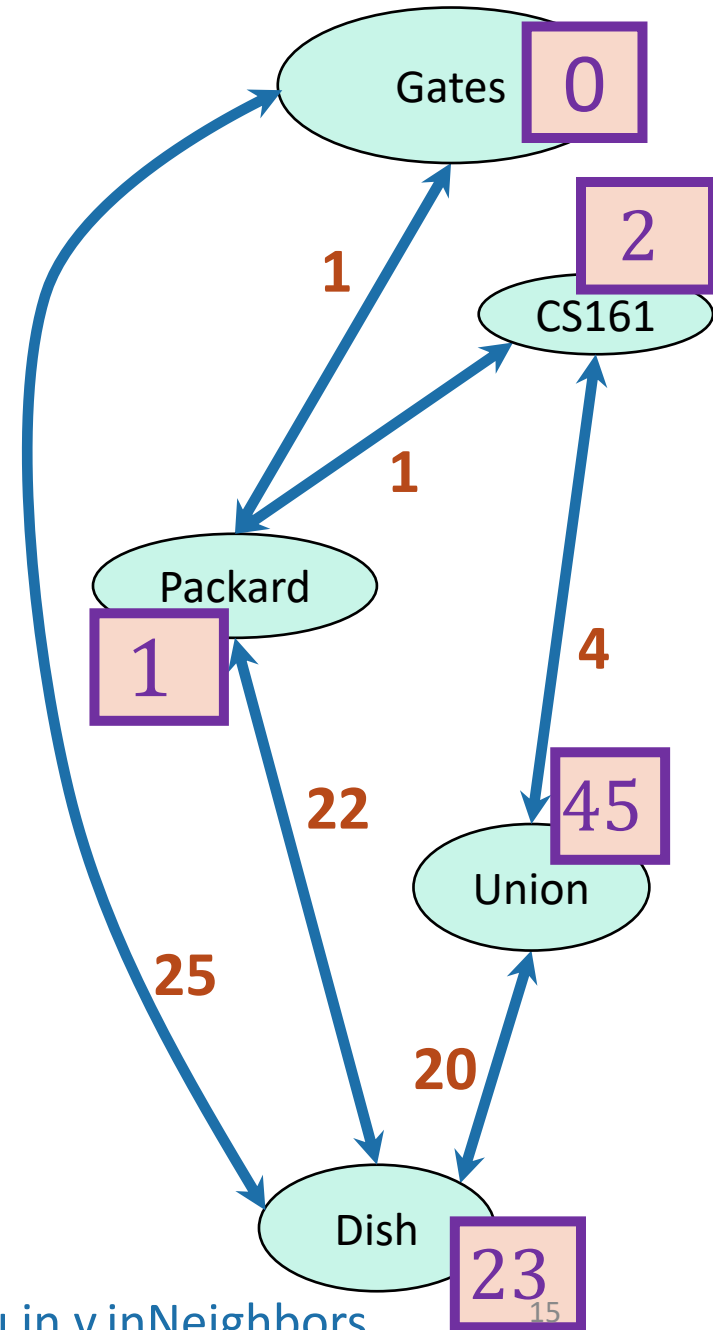


Bellman-Ford

How far is a node from Gates?

	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	25
$d^{(2)}$	0	1	2	45	23
$d^{(3)}$					
$d^{(4)}$					

- For $i=0, \dots, n-2$:
 - For v in V :
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v] , d^{(i)}[u] + w(u,v))$
 where we are also taking the min over all u in $v.inNeighbors$

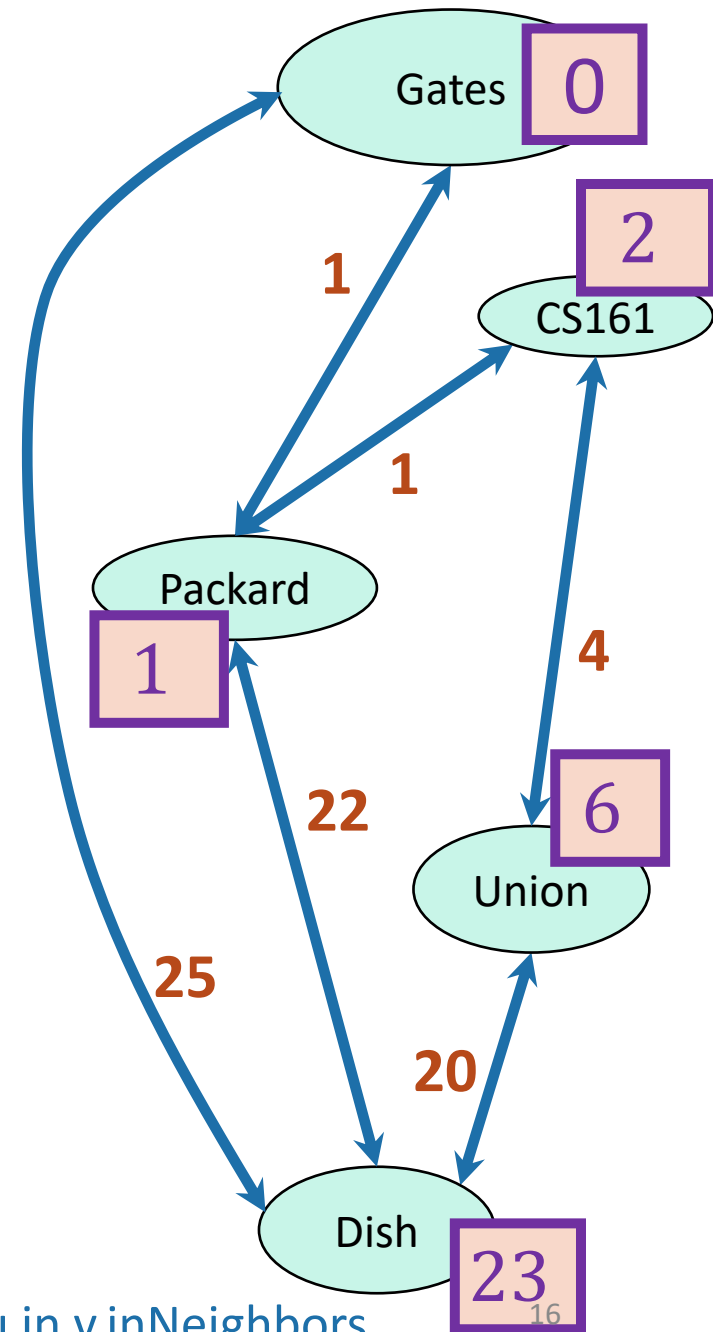


Bellman-Ford

How far is a node from Gates?

	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	25
$d^{(2)}$	0	1	2	45	23
$d^{(3)}$	0	1	2	6	23
$d^{(4)}$					

- For $i=0, \dots, n-2$:
 - For v in V :
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v] , d^{(i)}[u] + w(u,v))$
 where we are also taking the min over all u in $v.inNeighbors$



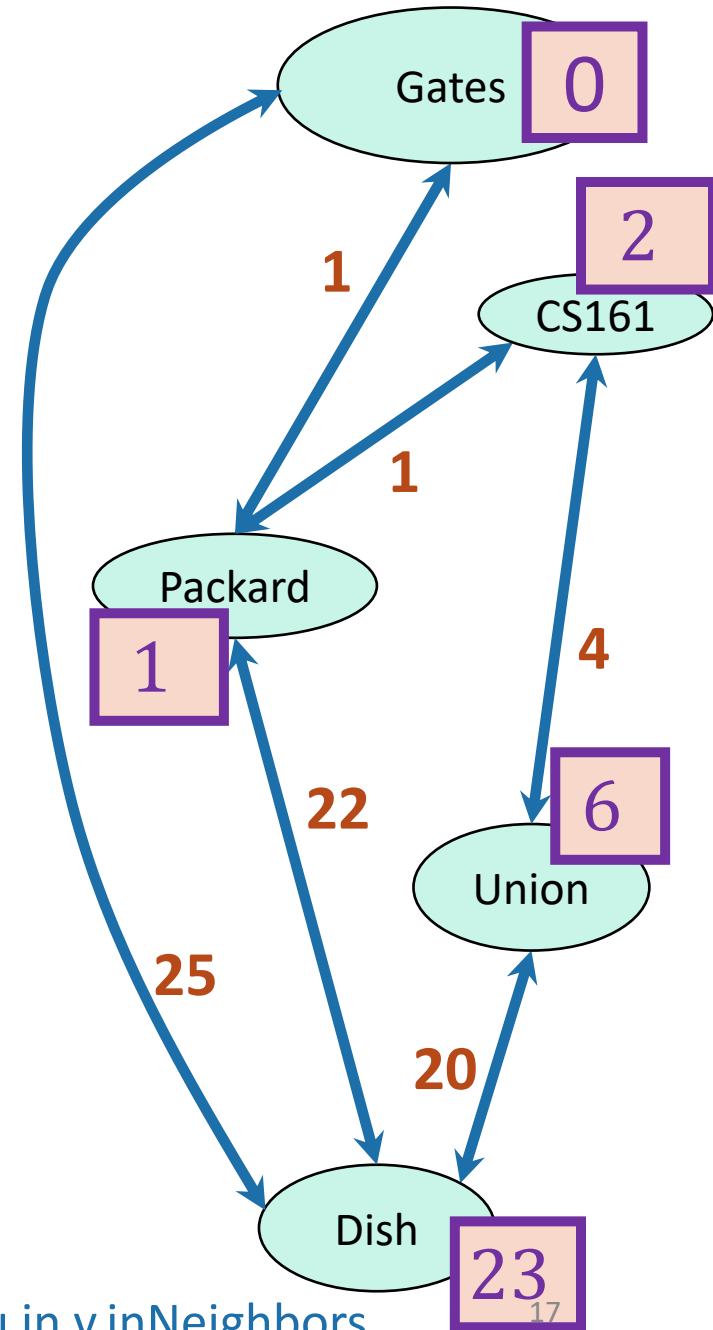
Bellman-Ford

How far is a node from Gates?

	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	25
$d^{(2)}$	0	1	2	45	23
$d^{(3)}$	0	1	2	6	23
$d^{(4)}$	0	1	2	6	23

These are the final distances!

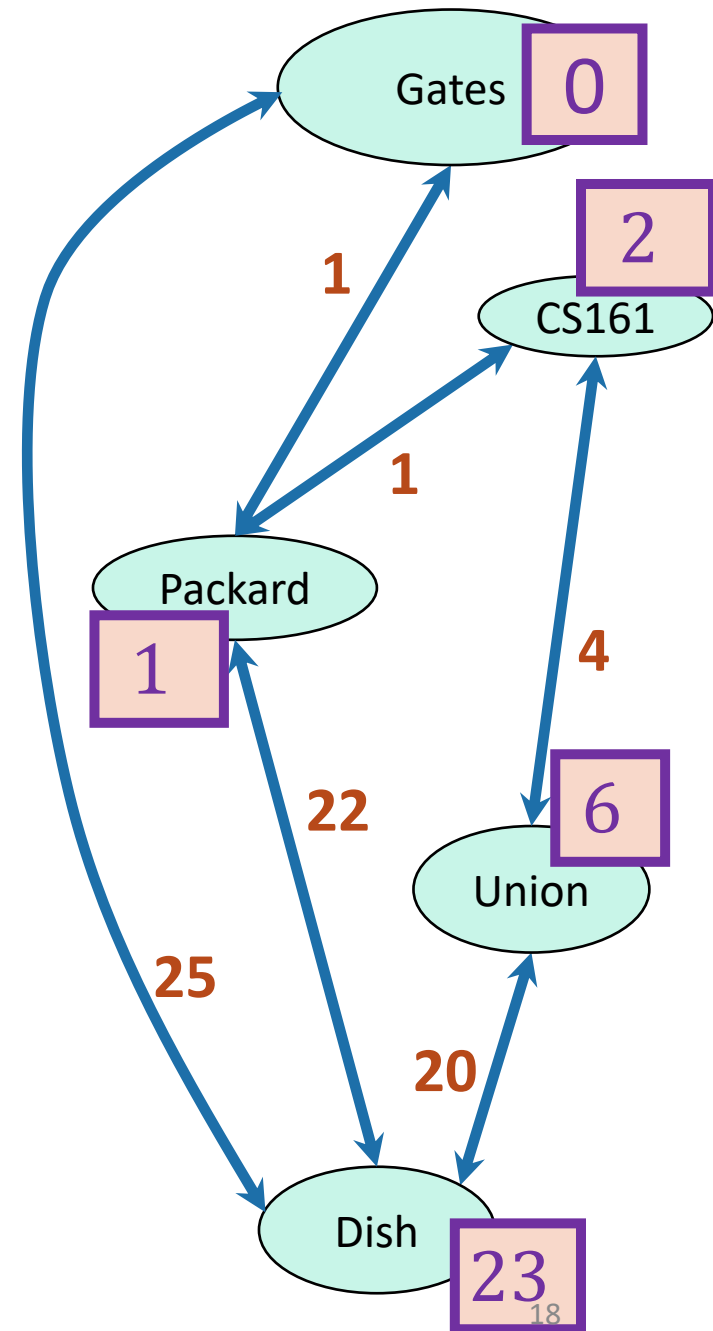
- For $i=0, \dots, n-2$:
 - For v in V :
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v] , d^{(i)}[u] + w(u,v))$
 where we are also taking the min over all u in $v.inNeighbors$



Interpretation of $d^{(i)}$

$d^{(i)}[v]$ is equal to the cost of the shortest path between s and v with at most i edges.

	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	25
$d^{(2)}$	0	1	2	45	23
$d^{(3)}$	0	1	2	6	23
$d^{(4)}$	0	1	2	6	23



Why does Bellman-Ford work?

- Inductive hypothesis:
 - $d^{(i)}[v]$ is equal to the cost of the shortest path between s and v **with at most i edges**.
- Conclusion:
 - $d^{(n-1)}[v]$ is equal to the cost of the shortest path between s and v **with at most $n-1$ edges**.


Do the base case and
inductive step!



Aside: simple paths

Assume there is no negative cycle.


- Then there is a shortest path from s to t , and moreover there is a **simple** shortest path.



This cycle isn't helping.
Just get rid of it.

- A **simple path** in a graph with n vertices has at most $n-1$ edges in it.

Can't add another edge
without making a cycle!



"Simple" means
that the path has
no cycles in it.



- So there is a shortest path with at most $n-1$ edges

Why does it work?

- Inductive hypothesis:
 - $d^{(i)}[v]$ is equal to the cost of the shortest path between s and v **with at most i edges**.
- Conclusion:
 - $d^{(n-1)}[v]$ is equal to the cost of the shortest path between s and v **with at most $n-1$ edges**.
 - **If there are no negative cycles**, $d^{(n-1)}[v]$ is equal to the cost of the shortest path.

Notice that negative edge **weights** are fine.
Just not negative cycles.

Bellman-Ford* algorithm

Bellman-Ford*(G,s):

- Initialize arrays $d^{(0)}, \dots, d^{(n-1)}$ of length n
- $d^{(0)}[v] = \infty$ for all v in V
- $d^{(0)}[s] = 0$
- **For** $i=0, \dots, n-2$:
 - **For** v in V :
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v] , \min_{u \text{ in } v.\text{inNbrs}} \{d^{(i)}[u] + w(u,v)\})$
- Now, $\text{dist}(s,v) = d^{(n-1)}[v]$ for all v in V .
 - (Assuming no negative cycles)

Here, Dijkstra picked a special vertex u and updated u 's neighbors – Bellman-Ford will update all the vertices.

*Slightly different than some versions of Bellman-Ford...but this way is pedagogically convenient for today's lecture.

Note on implementation

- Don't actually keep all n arrays around.
- Just keep two at a time: “last round” and “this round”

	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	25
$d^{(2)}$	0	1	2	45	23
$d^{(3)}$	0	1	2	6	23
$d^{(4)}$	0	1	2	6	23

Only need these two in order to compute $d^{(4)}$

Bellman-Ford take-aways

- Running time is $O(mn)$
 - For each of n rounds, update m edges.
- Works fine with negative edges.
- Does not work with negative cycles.
 - No algorithm can – shortest paths aren't defined if there are negative cycles.
- B-F can detect negative cycles!
 - See skipped slides to see how, or think about it on your own!


BF with negative cycles

SLIDE
SKIPPED
IN CLASS

	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	-3
$d^{(2)}$	0	-5	2	7	-3
$d^{(3)}$	-4	-5	-4	6	-3
$d^{(4)}$	-4	-5	-4	6	-7

This is not looking good!

- For $i=0, \dots, n-2$:
 - For v in V :
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v] , \min_{u \text{ in } v.\text{nbrs}} \{d^{(i)}[u] + w(u,v)\})$



BF with negative cycles

SLIDE
SKIPPED
IN CLASS


	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	-3
$d^{(2)}$	0	-5	2	7	-3
$d^{(3)}$	-4	-5	-4	6	-3
$d^{(4)}$	-4	-5	-4	6	-7

But **we can tell** that it's not looking good:

$d^{(5)}$	-4	-9	-4	3	-7
-----------	----	----	----	---	----

Some stuff changed!

- For $i=0, \dots, n-2$:
 - For v in V :
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v] , \min_{u \text{ in } v.\text{nbrs}} \{d^{(i)}[u] + w(u,v)\})$



Negative cycles in Bellman-Ford

- If there are no negative cycles:
 - Everything works as it should, and stabilizes in $n-1$ rounds.
- If there are negative cycles:
 - Not everything works as it should...
 - The $d[v]$ values will keep changing.
- Solution:
 - Go one round more and see if things change.

Bellman-Ford algorithm

Bellman-Ford*(G,s):

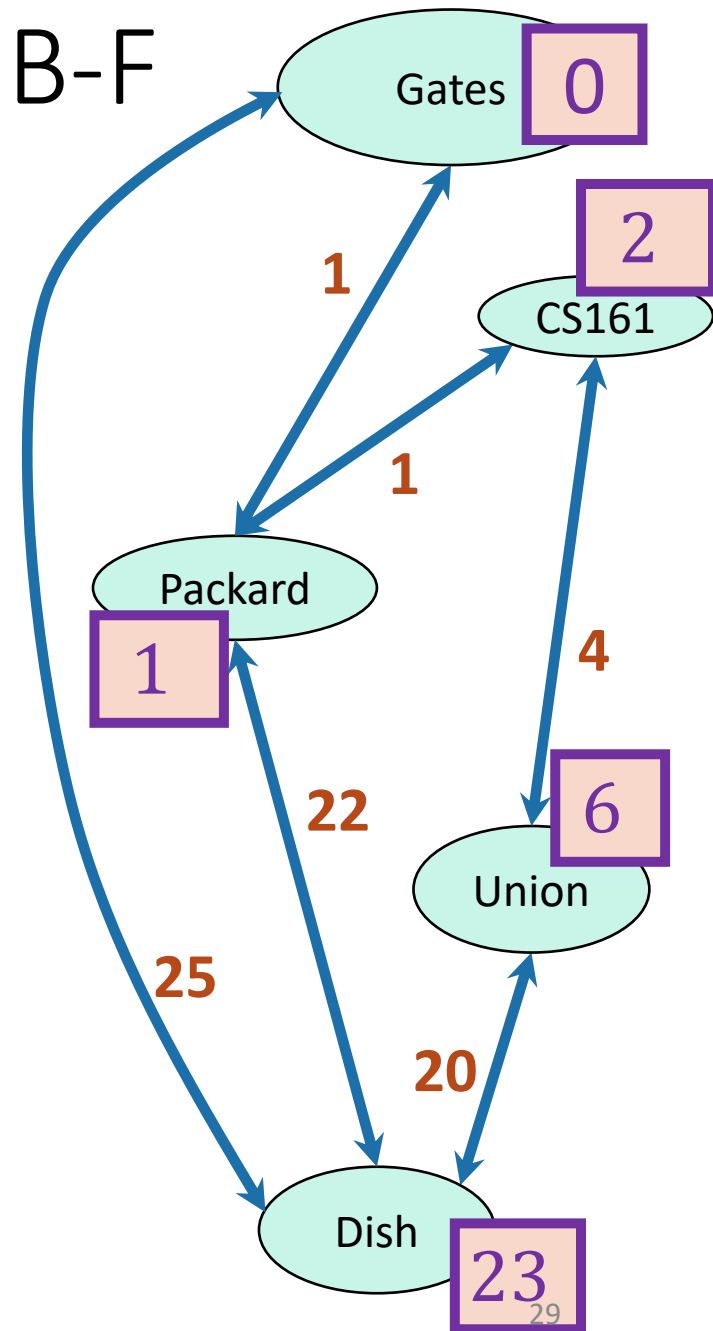
- $d^{(0)}[v] = \infty$ for all v in V
- $d^{(0)}[s] = 0$
- **For** $i=0, \dots, n-1$:
 - **For** v in V :
 - $d^{(i+1)}[v] \leftarrow \min(d^{(i)}[v] , \min_{u \text{ in } v.\text{inNeighbors}} \{d^{(i)}[u] + w(u,v)\})$
- **If** $d^{(n-1)} \neq d^{(n)}$:
 - **Return** **NEGATIVE CYCLE** 😞
- **Otherwise**, $\text{dist}(s,v) = d^{(n-1)}[v]$

Running time: $O(mn)$

Important thing about B-F for the rest of this lecture

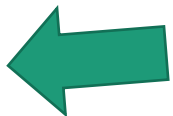
$d^{(i)}[v]$ is equal to the cost of the shortest path between s and v with at most i edges.

	Gates	Packard	CS161	Union	Dish
$d^{(0)}$	0	∞	∞	∞	∞
$d^{(1)}$	0	1	∞	∞	25
$d^{(2)}$	0	1	2	45	23
$d^{(3)}$	0	1	2	6	23
$d^{(4)}$	0	1	2	6	23



Bellman-Ford is an example of...
Dynamic Programming!

Today:

- Example of Dynamic programming: 
 - Fibonacci numbers.
 - (And Bellman-Ford)
- What is dynamic programming, exactly?
 - And why is it called “dynamic programming”?
- Another example: Floyd-Warshall algorithm
 - An “all-pairs” shortest path algorithm

Pre-Lecture exercise:

How not to compute Fibonacci Numbers


- Definition:
 - $F(n) = F(n-1) + F(n-2)$, with $F(1) = F(2) = 1$.
 - The first several are:
 - 1
 - 1
 - 2
 - 3
 - 5
 - 8
 - 13, 21, 34, 55, 89, 144,...
- Question:
 - Given n , what is $F(n)$?

Candidate algorithm

```
• def Fibonacci(n):  
    • if n == 0, return 0  
    • if n == 1, return 1  
    • return Fibonacci(n-1) + Fibonacci(n-2)
```

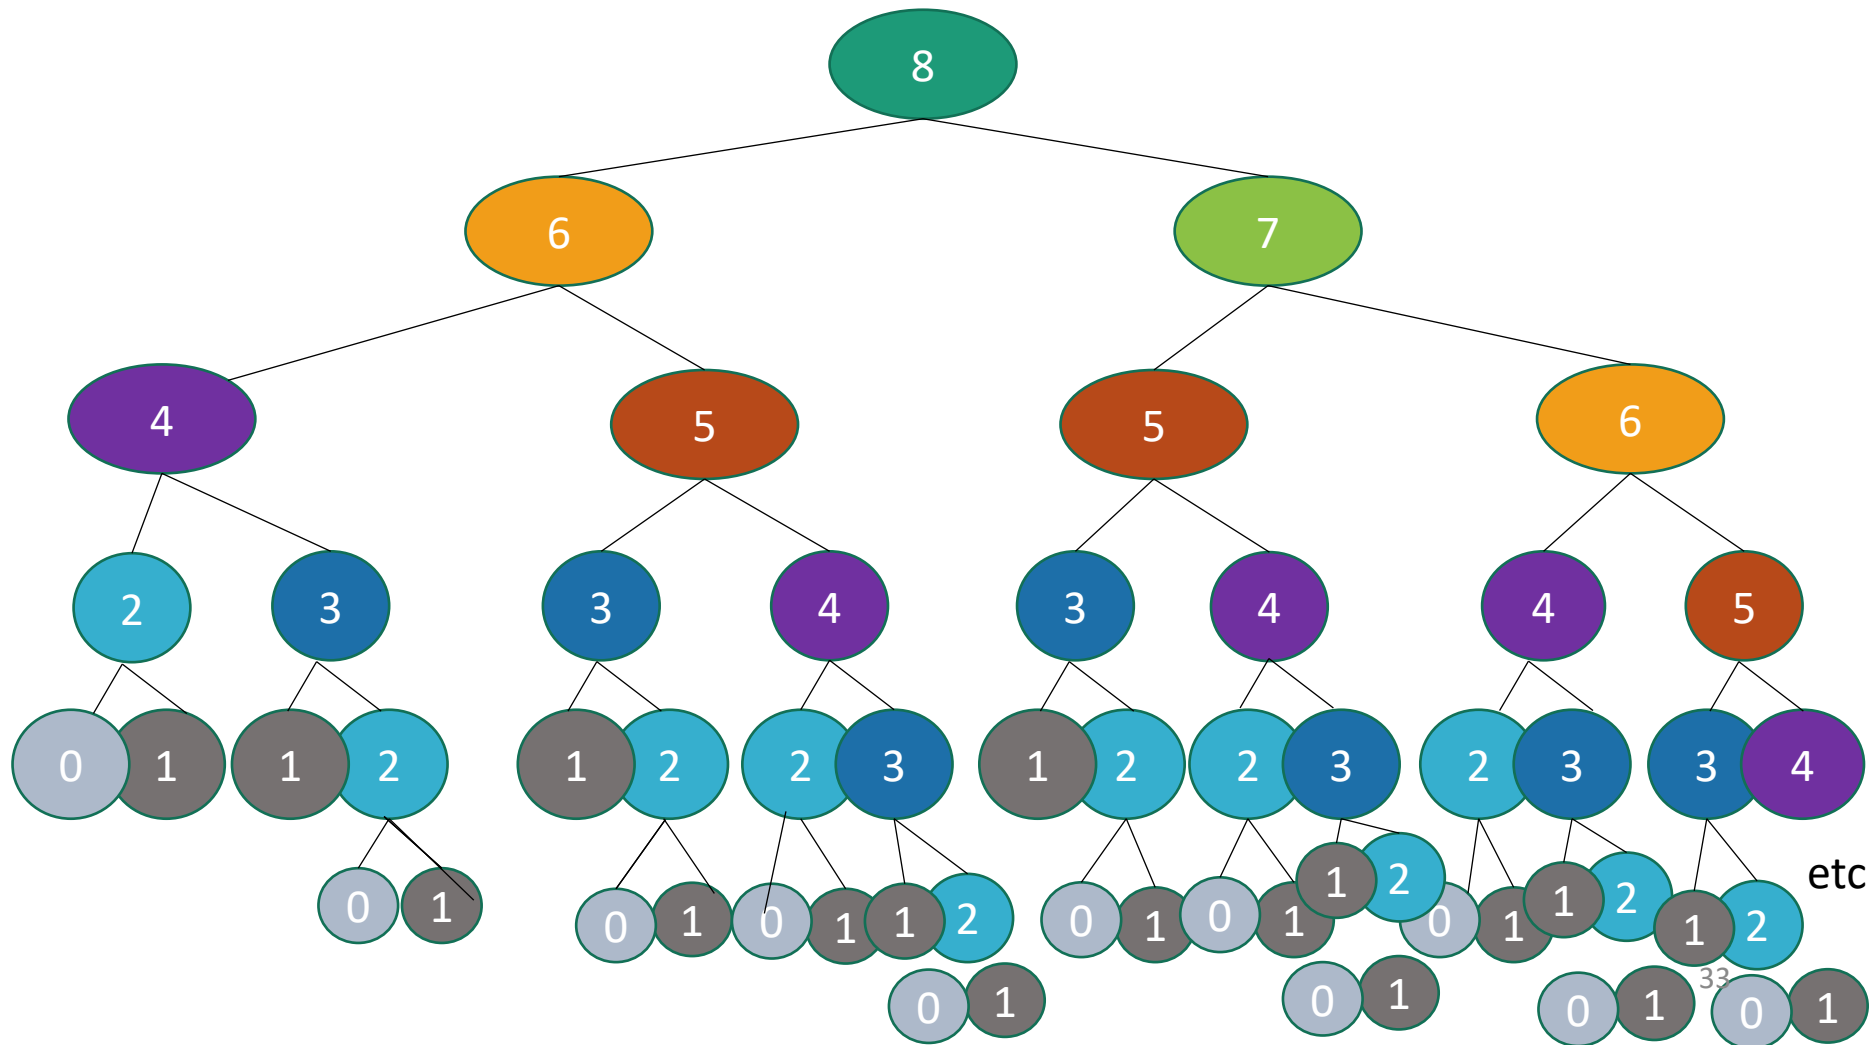
Running time?

- $T(n) = T(n-1) + T(n-2) + O(1)$
- $T(n) \geq T(n-1) + T(n-2)$ for $n \geq 2$
- So $T(n)$ grows *at least* as fast as the Fibonacci numbers themselves...
- You showed in HW1 that this is **EXPONENTIALLY QUICKLY!**




What's going on?
Consider $\text{Fib}(8)$

**That's a lot of
repeated
computation!**




Maybe this would be better:



```
def fasterFibonacci(n):  
    • F = [0, 1, None, None, ..., None ]  
      • \\ F has length n + 1  
    • for i = 2, ..., n:  
      • F[i] = F[i-1] + F[i-2]  
    • return F[n]
```

Much better running time!



This was an example of...

*Dynamic
programming!*

What is *dynamic programming*?

- It is an algorithm design paradigm
 - like divide-and-conquer is an algorithm design paradigm.
- Usually it is for solving **optimization problems**
 - eg, *shortest* path
 - (Fibonacci numbers aren't an optimization problem, but they are a good example of DP anyway...)

Elements of dynamic programming

1. Optimal sub-structure:

- Big problems break up into sub-problems.
 - Fibonacci: $F(i)$ for $i \leq n$
 - Bellman-Ford: Shortest paths with at most i edges for $i \leq n$
- The solution to a problem can be expressed in terms of solutions to smaller sub-problems.
 - Fibonacci:

$$F(i+1) = F(i) + F(i-1)$$

- Bellman-Ford:

$$d^{(i+1)}[v] \leftarrow \min\{d^{(i)}[v], \min_u \{d^{(i)}[u] + \text{weight}(u,v)\}\}$$

Shortest path with at most i edges from s to v

Shortest path with at most i edges from s to u .

Elements of dynamic programming

2. Overlapping sub-problems:


- The sub-problems overlap.
 - Fibonacci:
 - Both $F[i+1]$ and $F[i+2]$ directly use $F[i]$.
 - And lots of different $F[i+x]$ indirectly use $F[i]$.
 - Bellman-Ford:
 - Many different entries of $d^{(i+1)}$ will directly use $d^{(i)}[v]$.
 - And lots of different entries of $d^{(i+x)}$ will indirectly use $d^{(i)}[v]$.
- This means that we can save time by solving a sub-problem just once and storing the answer.

Elements of dynamic programming

- Optimal substructure.
 - Optimal solutions to sub-problems can be used to find the optimal solution of the original problem.
- Overlapping subproblems.
 - The subproblems show up again and again
- Using these properties, we can design a **dynamic programming** algorithm:
 - Keep a table of solutions to the smaller problems.
 - Use the solutions in the table to solve bigger problems.
 - At the end we can use information we collected along the way to find the solution to the whole thing.

Two ways to think about and/or implement DP algorithms

- Top down
- Bottom up




This picture isn't hugely relevant but I like it.

Larson

Bottom up approach

what we just saw.


- For Fibonacci:
- Solve the small problems first
 - fill in $F[0], F[1]$
- Then bigger problems
 - fill in $F[2]$
- ...
- Then bigger problems
 - fill in $F[n-1]$
- Then finally solve the real problem.
 - fill in $F[n]$



Bottom up approach


what we just saw.

- For Bellman-Ford:
- Solve the small problems first
 - fill in $d^{(0)}$
- Then bigger problems
 - fill in $d^{(1)}$
- ...
- Then bigger problems
 - fill in $d^{(n-2)}$
- Then finally solve the real problem.
 - fill in $d^{(n-1)}$



Top down approach


- Think of it like a recursive algorithm.
- To solve the big problem:
 - Recurse to solve smaller problems
 - Those recurse to solve smaller problems
 - etc..
- The difference from divide and conquer:
 - Keep track of what small problems you've already solved to prevent re-solving the same problem twice.
 - Aka, “**memo-ization**”



Example of top-down Fibonacci


- define a global list `F = [0,1,None, None, ..., None]`
- **def** `Fibonacci(n):`
 - **if** `F[n] != None:`
 - **return** `F[n]`
 - **else:**
 - `F[n] = Fibonacci(n-1) + Fibonacci(n-2)`
 - **return** `F[n]`

Memo-ization:
Keeps track (in F)
of the stuff you've
already done.



Memo-ization visualization

Collapse
repeated nodes
and don't do
the same work
twice!




Memo-ization Visualization

ctd

Collapse
repeated nodes
and don't do the
same work
twice!

But otherwise
treat it like the
same old
recursive
algorithm.

```
• define a global list F = [0,1,None, None, ..., None]
• def Fibonacci(n):
    • if F[n] != None:
        • return F[n]
    • else:
        • F[n] = Fibonacci(n-1) + Fibonacci(n-2)
    • return F[n]
```



What have we learned?

- ***Dynamic programming:***

- Paradigm in algorithm design.
- Uses **optimal substructure**
- Uses **overlapping subproblems**
- Can be implemented **bottom-up** or **top-down**.
- It's a fancy name for a pretty common-sense idea:



Why “*dynamic programming*” ?

- **Programming** refers to finding the optimal “program.”
 - as in, a shortest route is a *plan* aka a *program*.
- **Dynamic** refers to the fact that it’s multi-stage.
- But also it’s just a fancy-sounding name.



Manipulating computer code in an action movie?

Why “*dynamic programming*” ?


- Richard Bellman invented the name in the 1950's.
- At the time, he was working for the RAND Corporation, which was basically working for the Air Force, and government projects needed flashy names to get funded.
- From Bellman's autobiography:
 - “It's impossible to use the word, dynamic, in the pejorative sense...I thought dynamic programming was a good name. It was something not even a Congressman could object to.”

Floyd-Warshall Algorithm

Another example of DP

- This is an algorithm for **All-Pairs Shortest Paths (APSP)**
 - That is, I want to know the shortest path from u to v for **ALL pairs** u, v of vertices in the graph.
 - Not just from a special single source s .

Source	Destination				
	s	u	v	t	
	s	0	2	4	2
	u	1	0	2	0
	v	∞	∞	0	-2
	t	∞	∞	∞	0




Floyd-Warshall Algorithm

Another example of DP

- This is an algorithm for **All-Pairs Shortest Paths (APSP)**
 - That is, I want to know the shortest path from u to v for **ALL pairs** u, v of vertices in the graph.
 - Not just from a special single source s .
- Naïve solution (if we want to handle negative edge weights):
 - For all s in G :
 - Run Bellman-Ford on G starting at s .
 - Time $O(n \cdot nm) = O(n^2m)$,
 - may be as bad as n^4 if $m=n^2$

Can we do better?

Optimal substructure



Optimal substructure

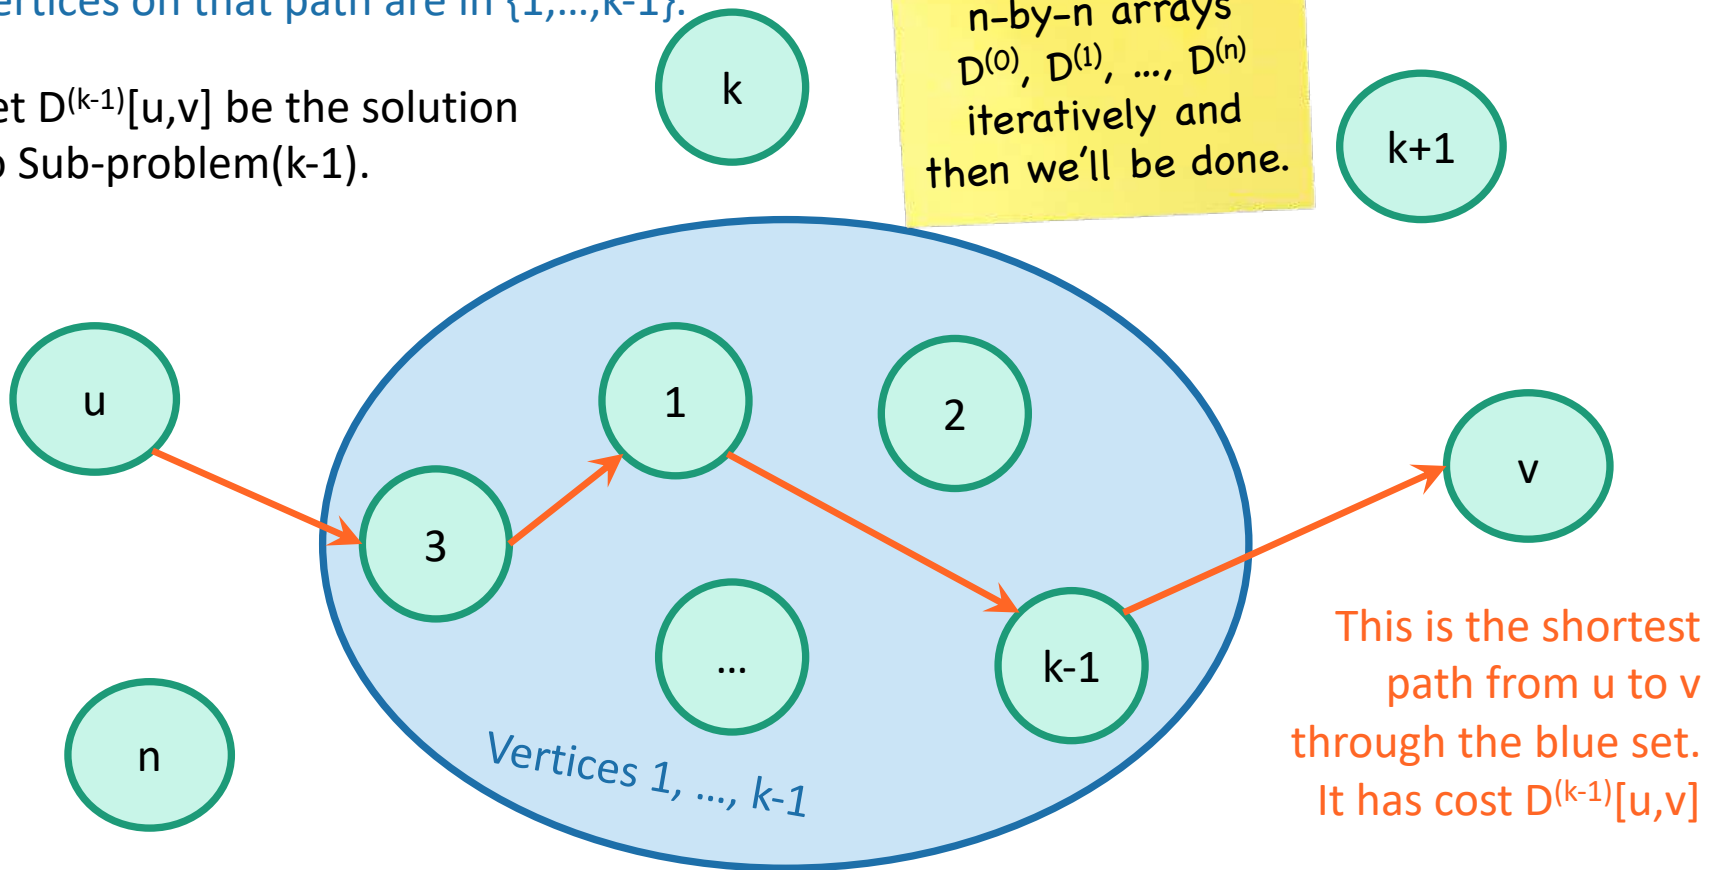
Label the vertices $1, 2, \dots, n$
(We omit some edges in the picture below – meant to be a cartoon, not an example).

Sub-problem($k-1$):

For all pairs, u, v , find the cost of the shortest path from u to v , so that all the internal vertices on that path are in $\{1, \dots, k-1\}$.

Let $D^{(k-1)}[u, v]$ be the solution to Sub-problem($k-1$).

Our DP algorithm will fill in the n -by- n arrays $D^{(0)}, D^{(1)}, \dots, D^{(n)}$ iteratively and then we'll be done.



Optimal substructure

Label the vertices $1, 2, \dots, n$
(We omit some edges in the picture below – meant to be a cartoon, not an example).


Sub-problem($k-1$):

For all pairs, u, v , find the cost of the shortest path from u to v , so that all the internal vertices on that path are in $\{1, \dots, k-1\}$.

Let $D^{(k-1)}[u, v]$ be the solution to Sub-problem($k-1$).


Our DP algorithm will fill in the n -by- n arrays $D^{(0)}, D^{(1)}, \dots, D^{(n)}$ iteratively and then we'll be done.

Question: How can we find $D^{(k)}[u, v]$ using $D^{(k-1)}$?



How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

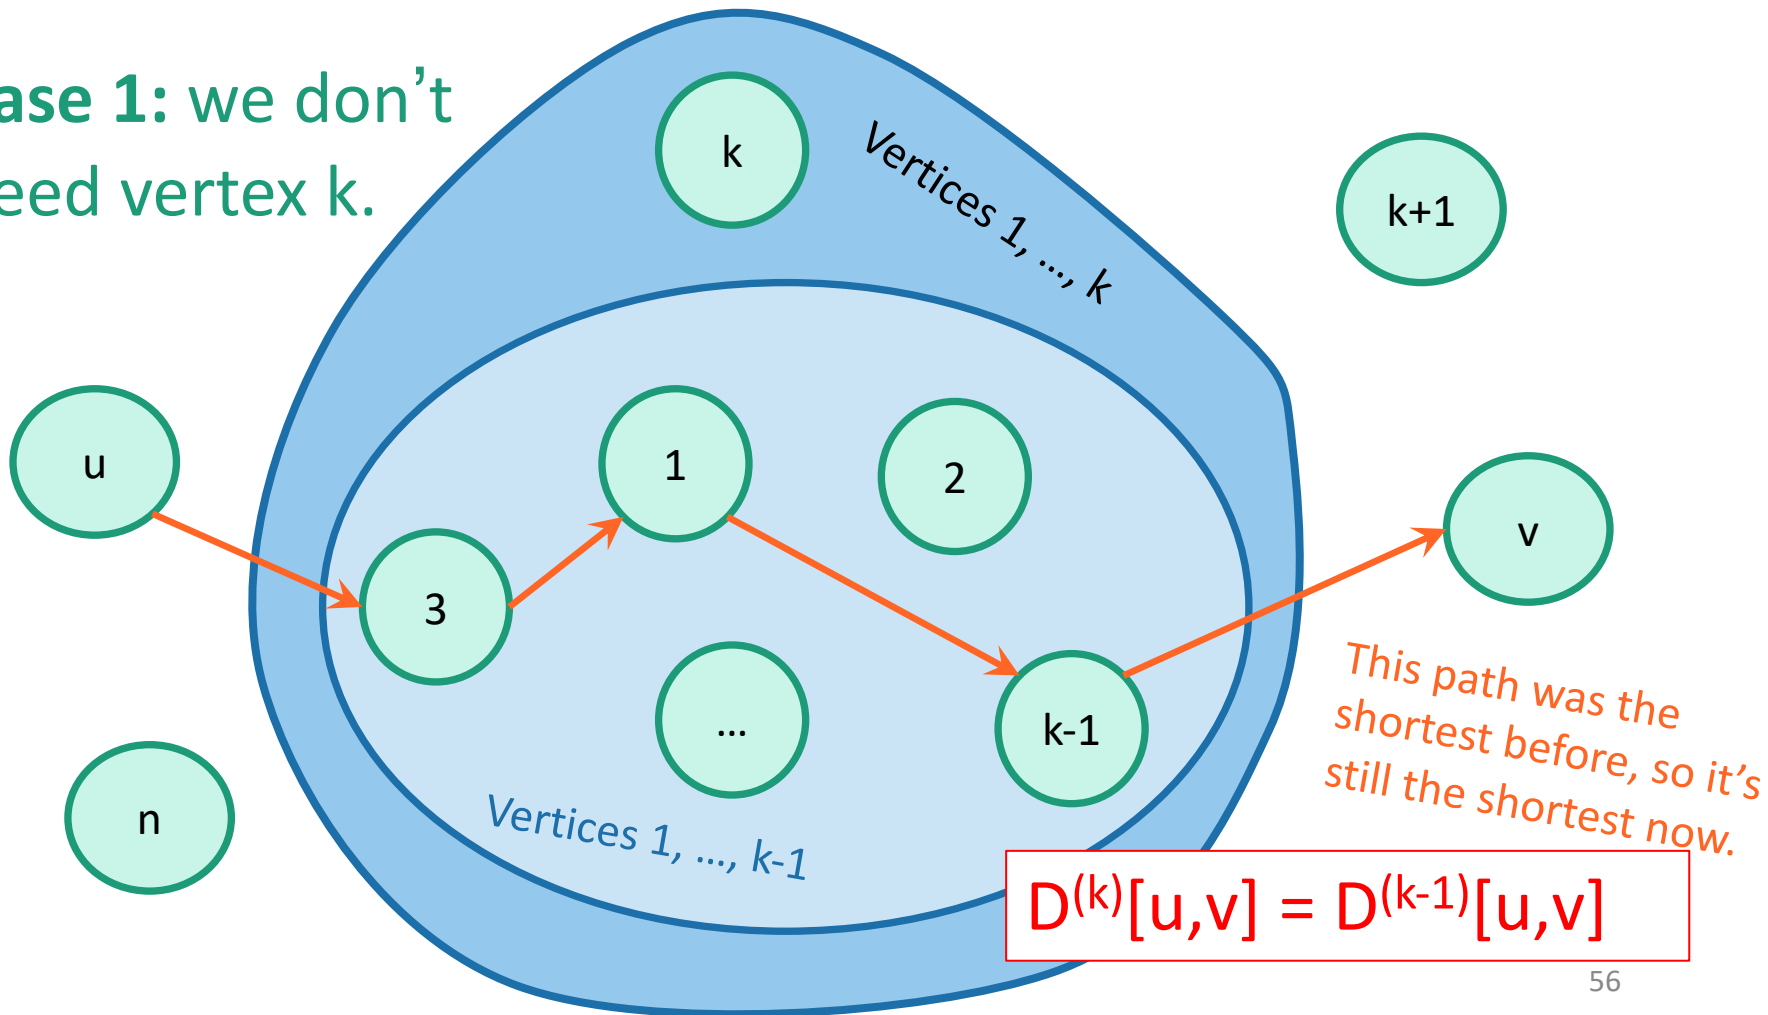
$D^{(k)}[u,v]$ is the cost of the shortest path from u to v so that all internal vertices on that path are in $\{1, \dots, k\}$.



How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

$D^{(k)}[u,v]$ is the cost of the shortest path from u to v so that all internal vertices on that path are in $\{1, \dots, k\}$.


Case 1: we don't need vertex k .





How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

$D^{(k)}[u,v]$ is the cost of the shortest path from u to v so that all internal vertices on that path are in $\{1, \dots, k\}$.


Case 2: we need vertex k .



Case 2 continued

- Suppose there are **no negative cycles**.
 - Then WLOG the shortest path from u to v through $\{1, \dots, k\}$ is **simple**.
- If that path passes through k , it must look like this: 
- This path is the shortest path from u to k through $\{1, \dots, k-1\}$.
 - sub-paths of shortest paths are shortest paths
- Similarly for this path. 


Case 2: we need vertex k .



$$D^{(k)}[u, v] = D^{(k-1)}[u, k] + D^{(k-1)}[k, v]_{58}$$


How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

Case 1: we don't need vertex k .



$$D^{(k)}[u,v] = D^{(k-1)}[u,v]$$

Case 2: we need vertex k .



$$D^{(k)}[u,v] = D^{(k-1)}[u,k] + D^{(k-1)}[k,v]$$

How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

- $D^{(k)}[u,v] = \min\{ D^{(k-1)}[u,v], D^{(k-1)}[u,k] + D^{(k-1)}[k,v] \}$

Case 1: Cost of
shortest path
through $\{1, \dots, k-1\}$

Case 2: Cost of shortest path
from u to k and then from k to v
through $\{1, \dots, k-1\}$

- Optimal substructure:
 - We can solve the big problem using solutions to smaller problems.
- Overlapping sub-problems:
 - $D^{(k-1)}[k,v]$ can be used to help compute $D^{(k)}[u,v]$ for lots of different u 's.

How can we find $D^{(k)}[u,v]$ using $D^{(k-1)}$?

- $D^{(k)}[u,v] = \min\{ D^{(k-1)}[u,v], D^{(k-1)}[u,k] + D^{(k-1)}[k,v] \}$

Case 1: Cost of
shortest path
through $\{1, \dots, k-1\}$

Case 2: Cost of shortest path
from **u** to **k** and then from **k** to **v**
through $\{1, \dots, k-1\}$

- Using our *Dynamic programming* paradigm, this immediately gives us an algorithm!



Floyd-Warshall algorithm

- Initialize n-by-n arrays $D^{(k)}$ for $k = 0, \dots, n$

- $D^{(k)}[u,u] = 0$ for all u , for all k
- $D^{(k)}[u,v] = \infty$ for all $u \neq v$, for all k
- $D^{(0)}[u,v] = \text{weight}(u,v)$ for all (u,v) in E .

The base case checks out: the only path through zero other vertices are edges directly from u to v .

- **For** $k = 1, \dots, n$:

- **For** pairs u,v in V^2 :

- $D^{(k)}[u,v] = \min\{ D^{(k-1)}[u,v], D^{(k-1)}[u,k] + D^{(k-1)}[k,v] \}$

- **Return** $D^{(n)}$

This is a bottom-up *Dynamic programming* algorithm.


We've basically just shown

- Theorem:

If there are **no negative cycles** in a weighted directed graph G , then the Floyd-Warshall algorithm, running on G , returns a matrix $D^{(n)}$ so that:

$$D^{(n)}[u,v] = \text{distance between } u \text{ and } v \text{ in } G.$$

Work out the
details of a proof!



- Running time: $O(n^3)$

- Better than running Bellman-Ford n times!

- Storage:

- Need to store **two** n -by- n arrays, and the original graph.

As with Bellman-Ford, we don't really need to store all n of the $D^{(k)}$.

What if there *are* negative cycles?

- Just like Bellman-Ford, Floyd-Warshall can detect negative cycles:
 - “Negative cycle” means that there’s some v so that there is a path from v to v that has cost < 0 .
 - Aka, $D^{(n)}[v,v] < 0$.
- Algorithm:
 - Run Floyd-Warshall as before.
 - If there is some v so that $D^{(n)}[v,v] < 0$:
 - **return negative cycle.**

What have we learned?

- The Floyd-Warshall algorithm is another example of *dynamic programming*.
- It computes All Pairs Shortest Paths in a directed weighted graph in time $O(n^3)$.

Can we do better than $O(n^3)$?

Nothing on this slide is required knowledge for this class

- There is an algorithm that runs in time $O(n^3/\log^{100}(n))$.
 - *[Williams, “Faster APSP via Circuit Complexity”, STOC 2014]*
- If you can come up with an algorithm for All-Pairs-Shortest-Path that runs in time $O(n^{2.99})$, that would be a really big deal.
 - Let me know if you can!
 - See *[Abboud, Vassilevska-Williams, “Popular conjectures imply strong lower bounds for dynamic problems”, FOCS 2014]* for some evidence that this is a very difficult problem!

Recap

- Two shortest-path algorithms:
 - Bellman-Ford for single-source shortest path
 - Floyd-Warshall for all-pairs shortest path
- ***Dynamic programming!***
 - This is a fancy name for:
 - Break up an optimization problem into smaller problems
 - The optimal solutions to the sub-problems should be sub-solutions to the original problem.
 - Build the optimal solution iteratively by filling in a table of sub-solutions.
 - Take advantage of overlapping sub-problems!

Next time

- More examples of *dynamic programming*!

We will stop bullets with our
action-packed coding skills,
and also maybe find longest
common subsequences.



- No pre-lecture exercise for next time: go over your exam instead!